




# Apache Kafka Kibana Data pipeline














Somkiat


Home









Update Info 1
View Activity Log 10+
...


Timeline
About
Friends 3,138
Photos
More ▾



When did you work at Opendream?
×





...
22 Pending Items



Intro


Software Craftsmanship


Software Practitioner at สยามชำนาญกิจ พ.ศ. 2556


Agile Practitioner and Technical at SPRINT3r


Post

Photo/Video

Live Video

Life Event


What's on your mind?


Public ▾

Post



**Somkiat Puisungnoen**
15 mins · Bangkok · ▾

Java and Bigdata

...



Facebook interface for the page **somkiat.cc**. The top navigation bar includes the Facebook logo, the page name, a search bar, and icons for Home, Messages, Notifications (3), Insights, Publishing Tools, Settings, and Help.

The main content area features a large video player showing a man in a white Superman t-shirt with "SOMKIAT.CC" printed on it, posing against a white wall. A blue call-to-action button is overlaid on the video: "Help people take action on this Page. x". Below the video are buttons for "Liked", "Following", "Share", and a menu icon.

The left sidebar contains the page name **somkiat.cc**, the handle **@somkiat.cc**, and a menu with options: Home, Posts, Videos, and Photos.

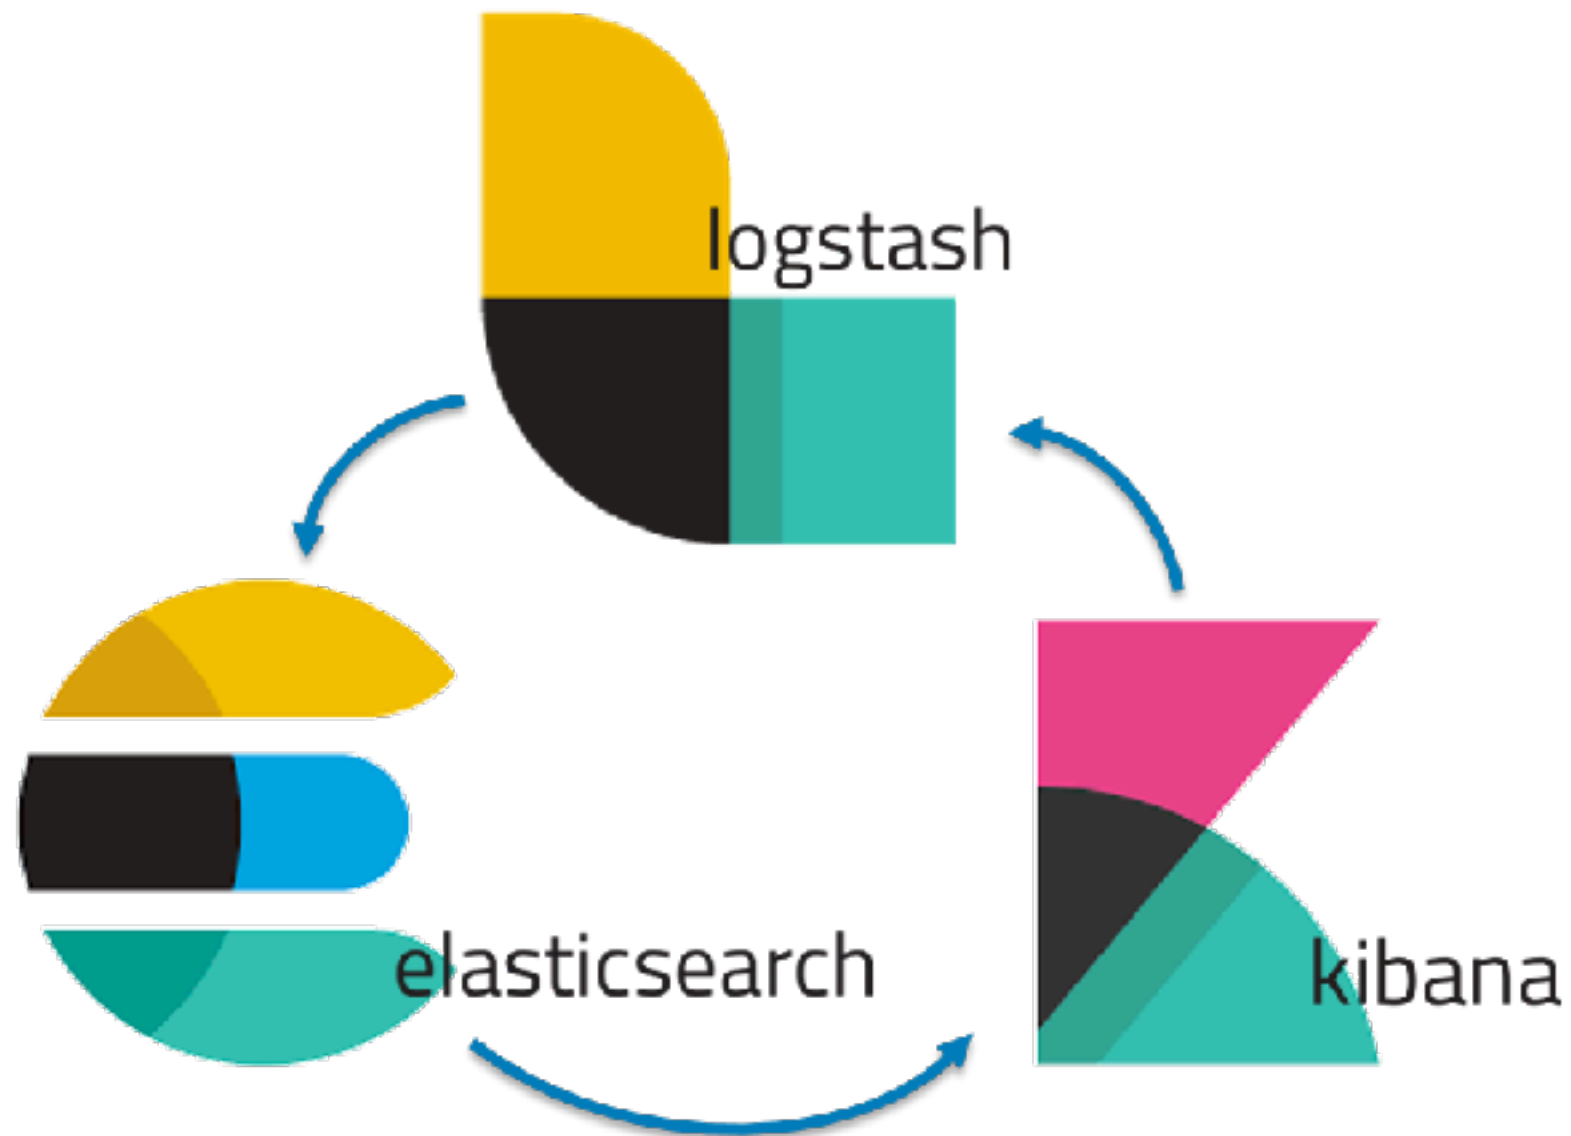


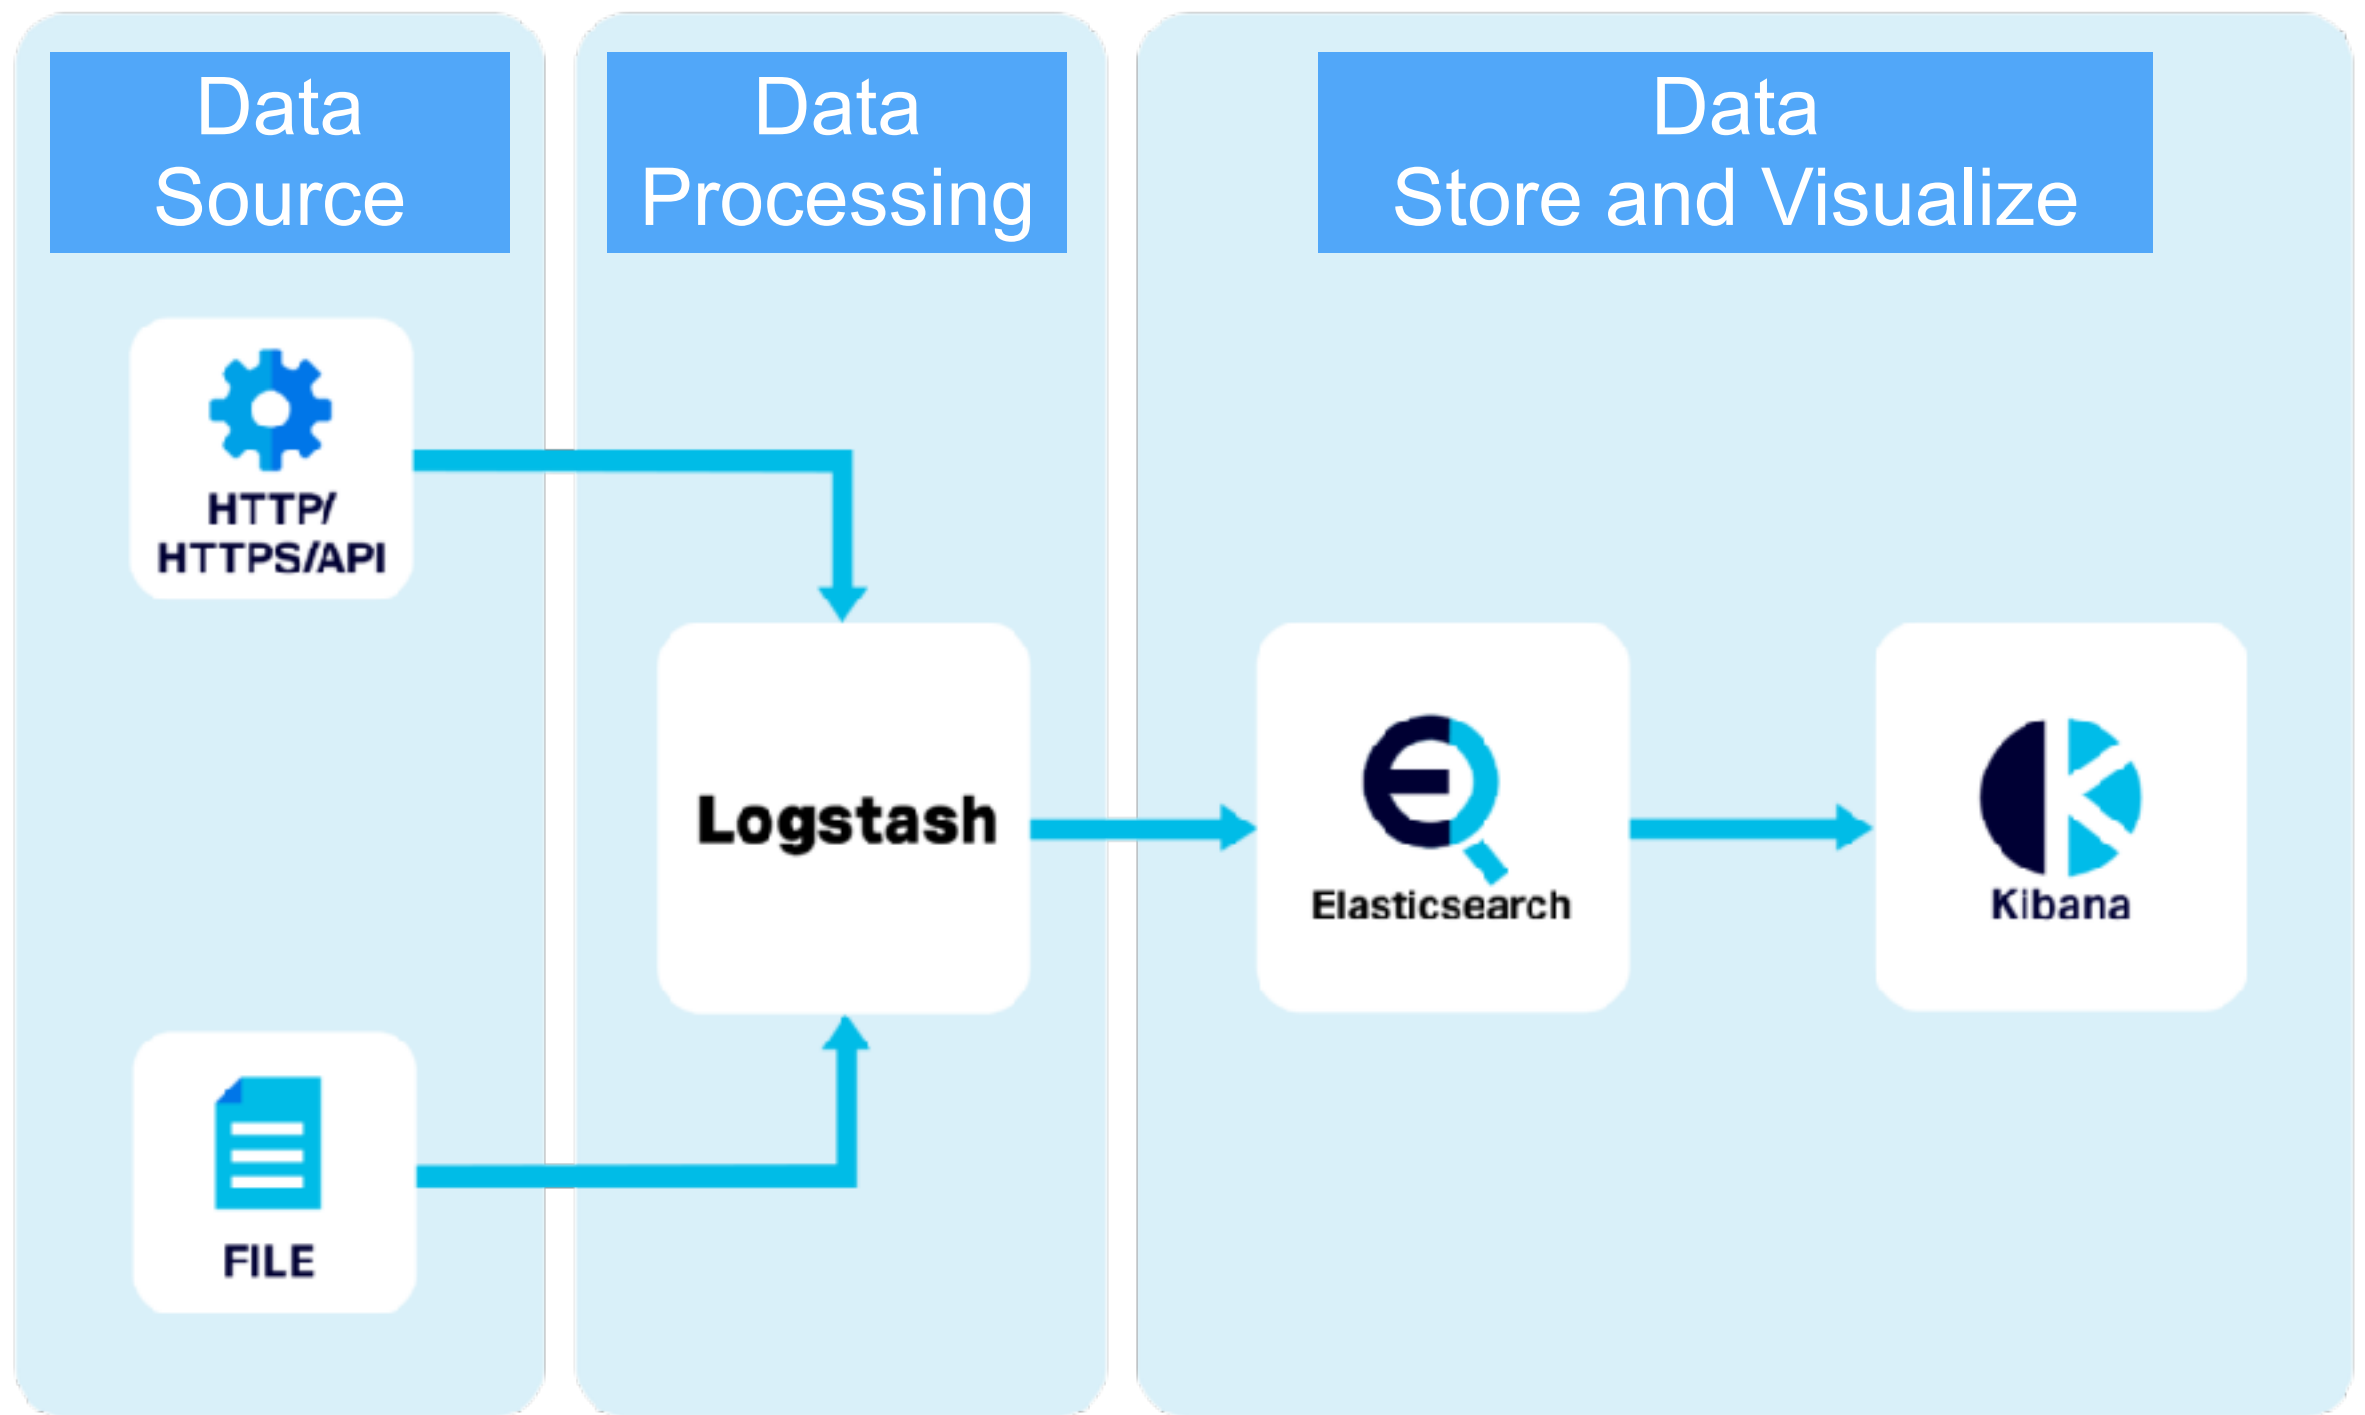
# Agenda

- Recap ELK stack
- Apache Kafka
- Working with Kibana
- Data pipeline
- Demo



# ELK stack





Transform data

Logstash

Data Store

Elasticsearch

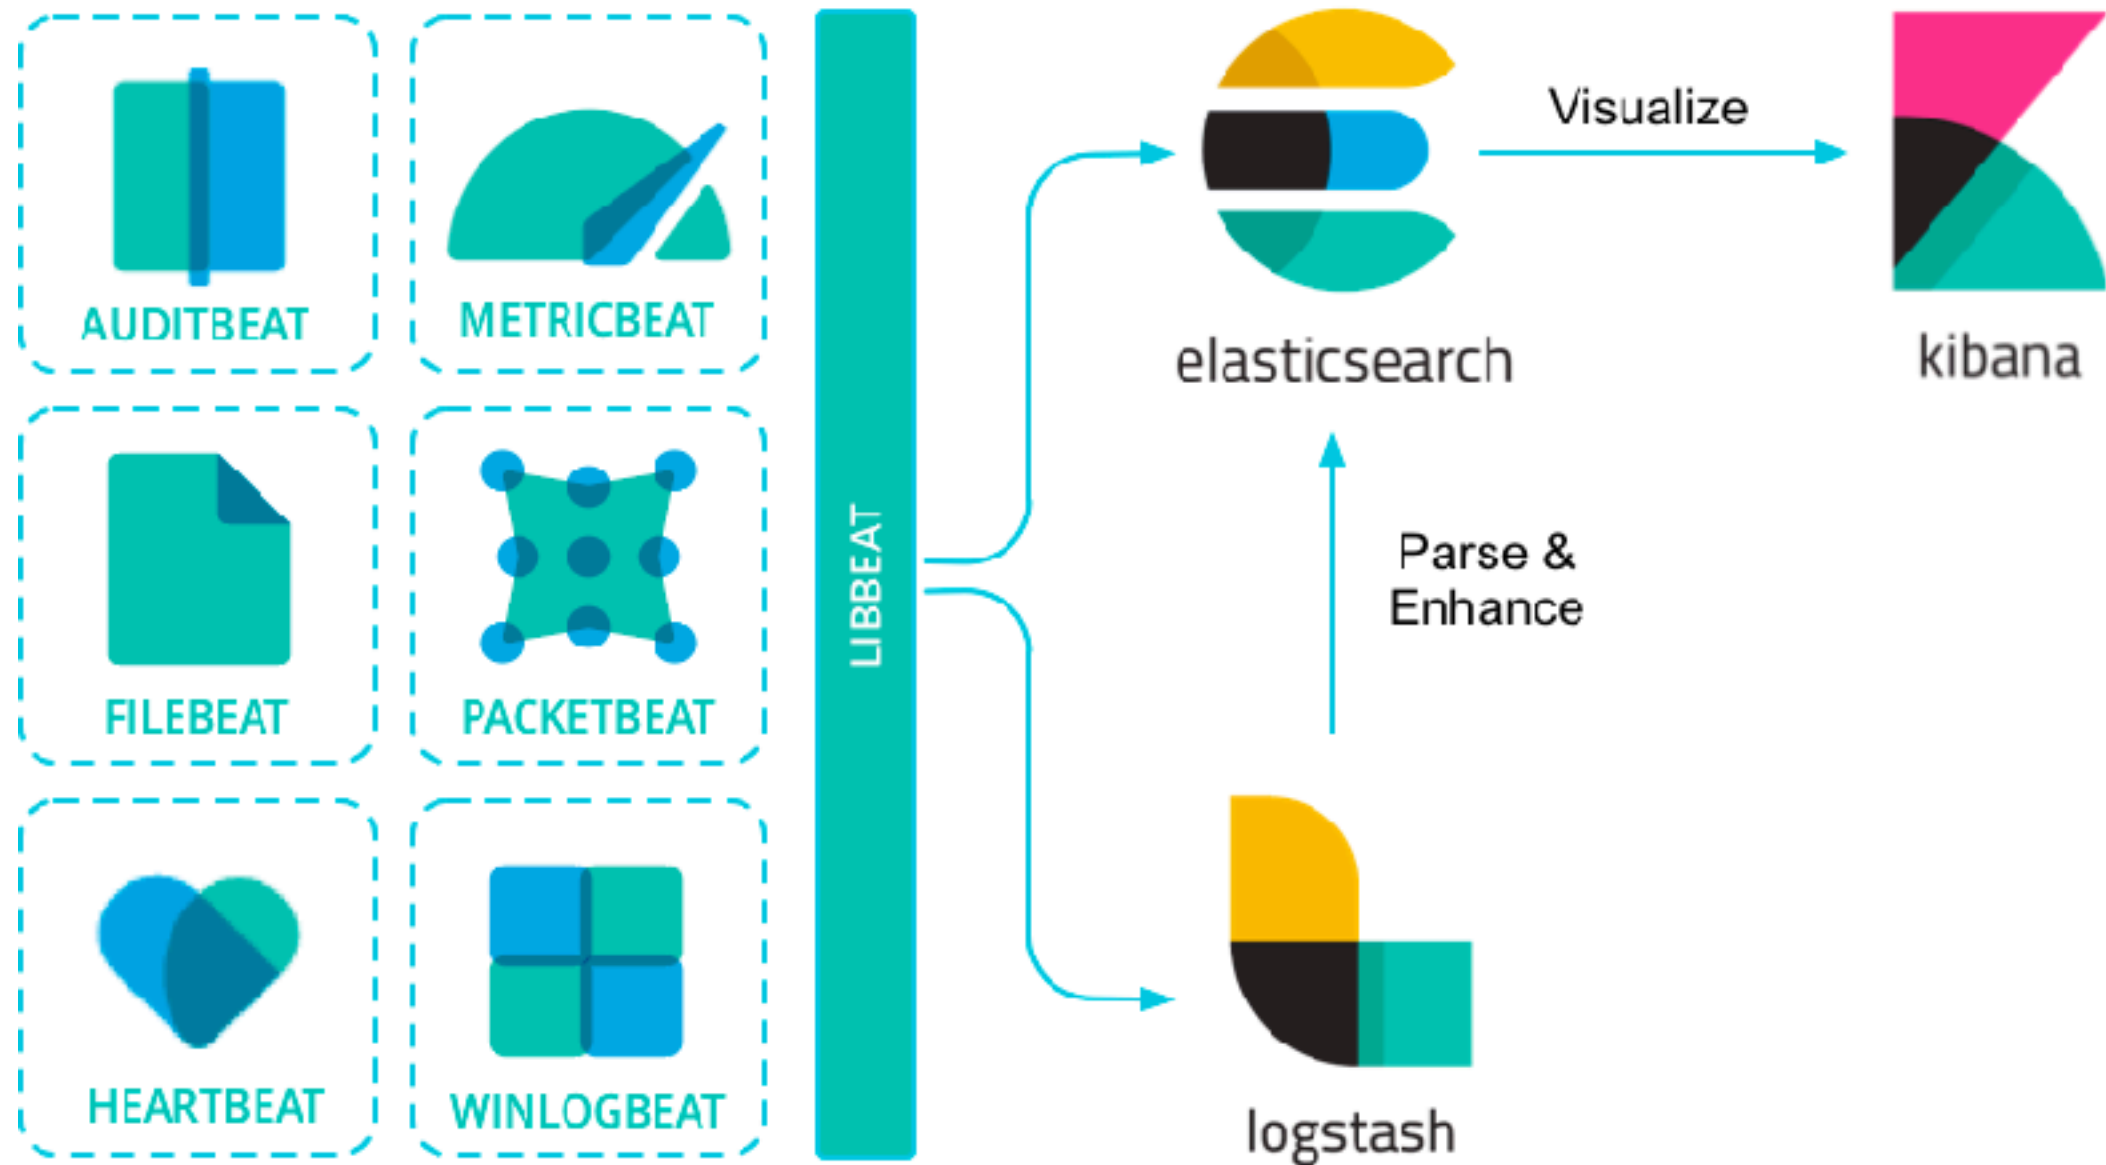
Visualize

Kibana





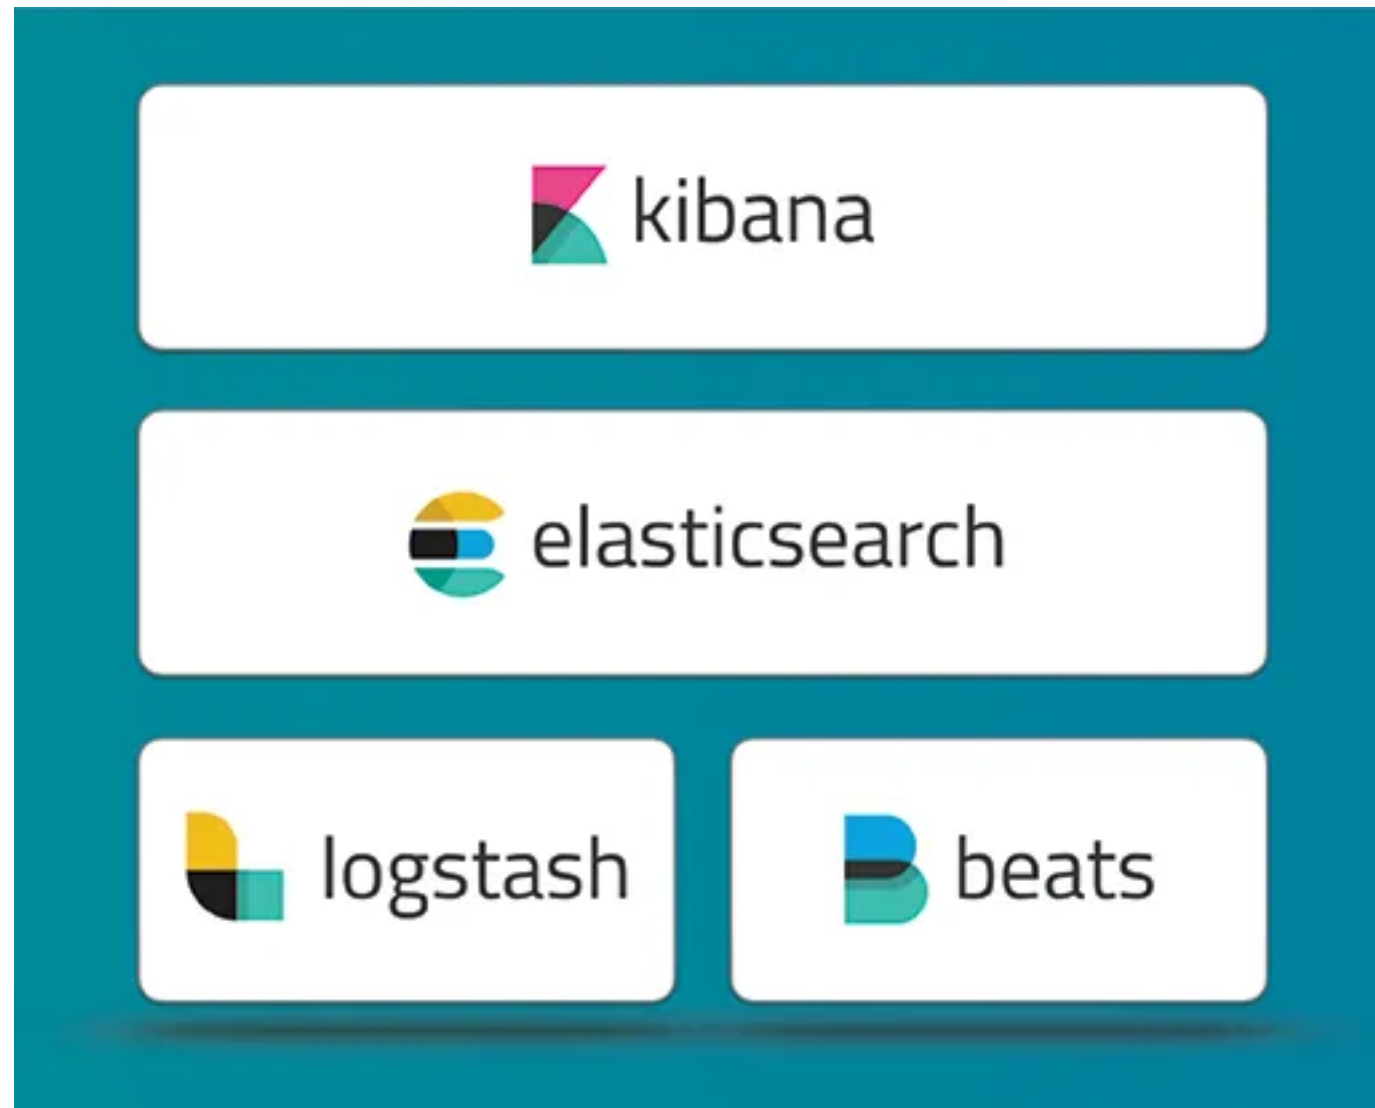
# Beat



<https://www.elastic.co/guide/en/beats/libbeat/current/index.html>



# ELK stack



# Beat

| Purpose            | Library      |
|--------------------|--------------|
| Audit data         | Auditbeat    |
| Log files          | Filebeat     |
| Cloud data         | Functionbeat |
| Availability       | Heartbeat    |
| Metrics            | Metricbeat   |
| Network traffic    | Packetbeat   |
| Windows event logs | Winlogbeat   |

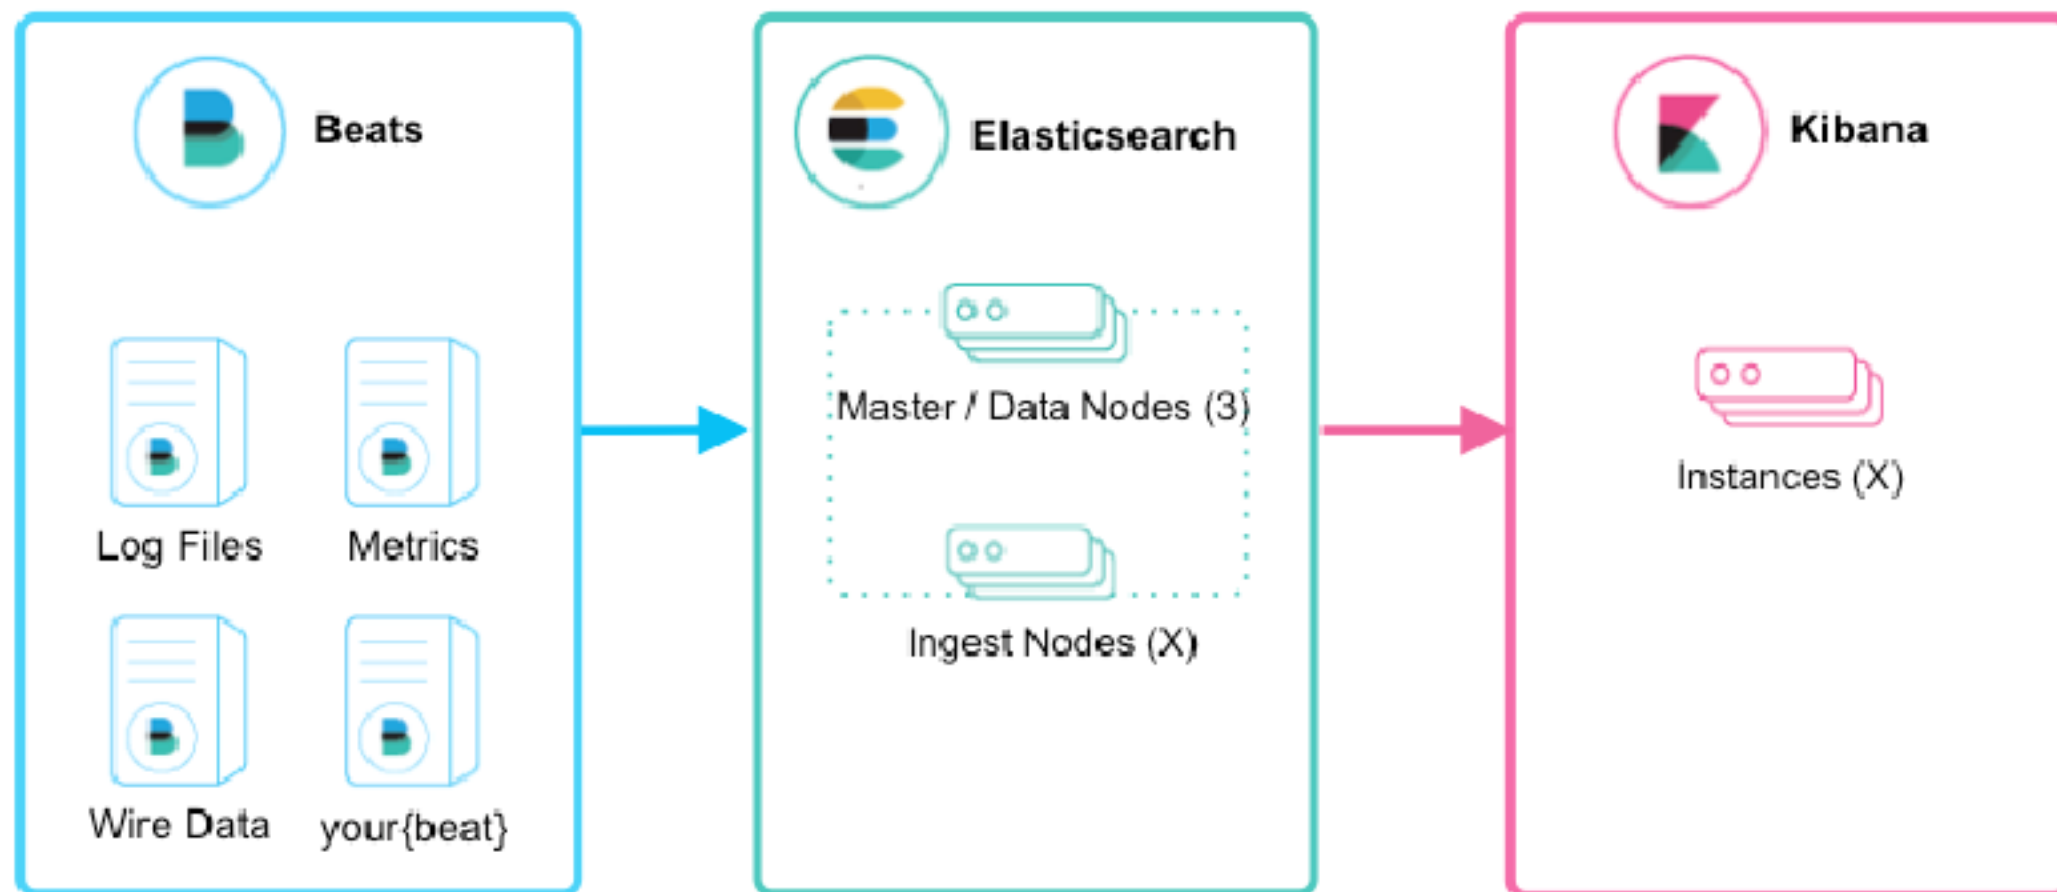
<https://www.elastic.co/guide/en/beats/libbeat/current/beats-reference.html>



# Scaling ELK Stack



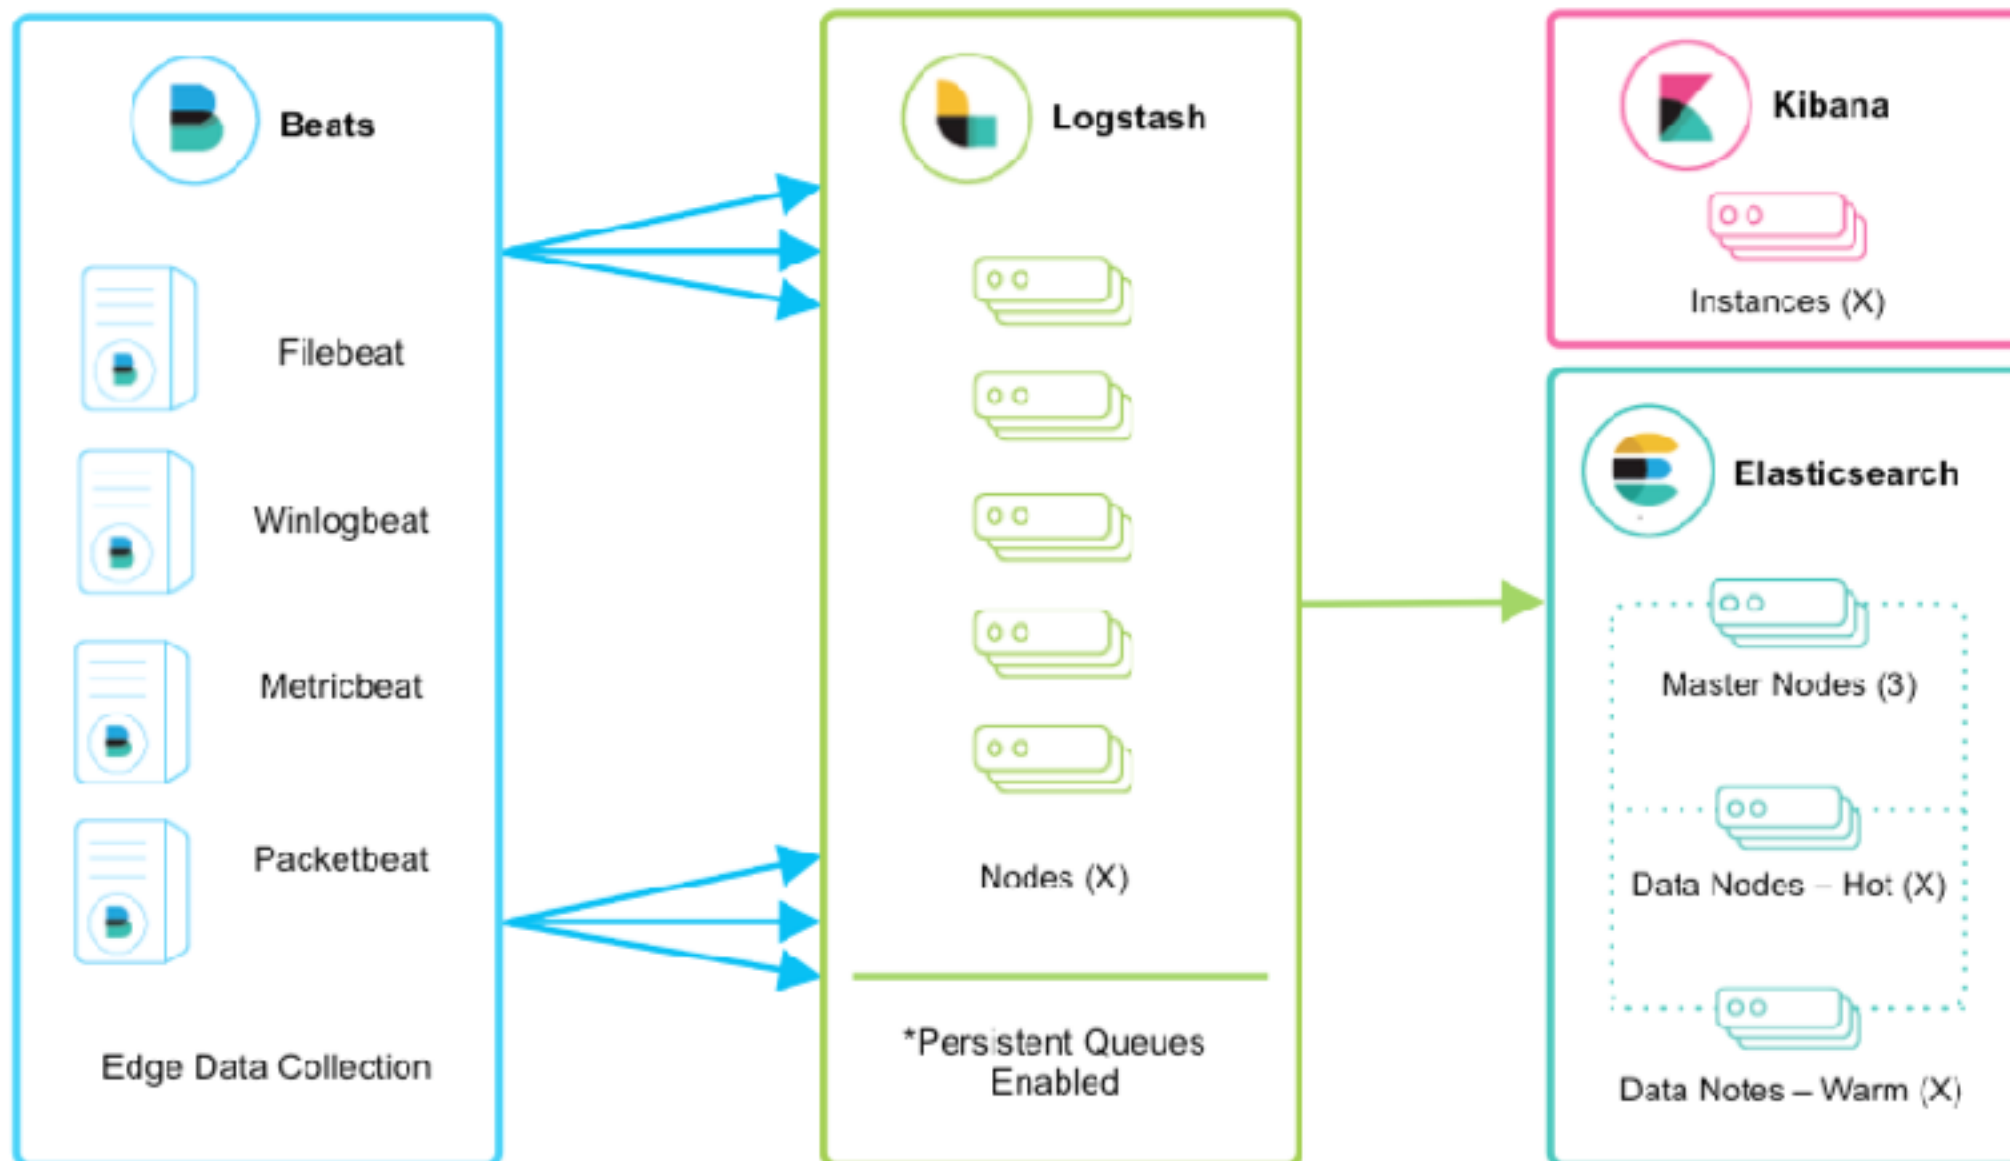
# Basic



<https://www.elastic.co/docs/reference/logstash/deploying-scaling-logstash>



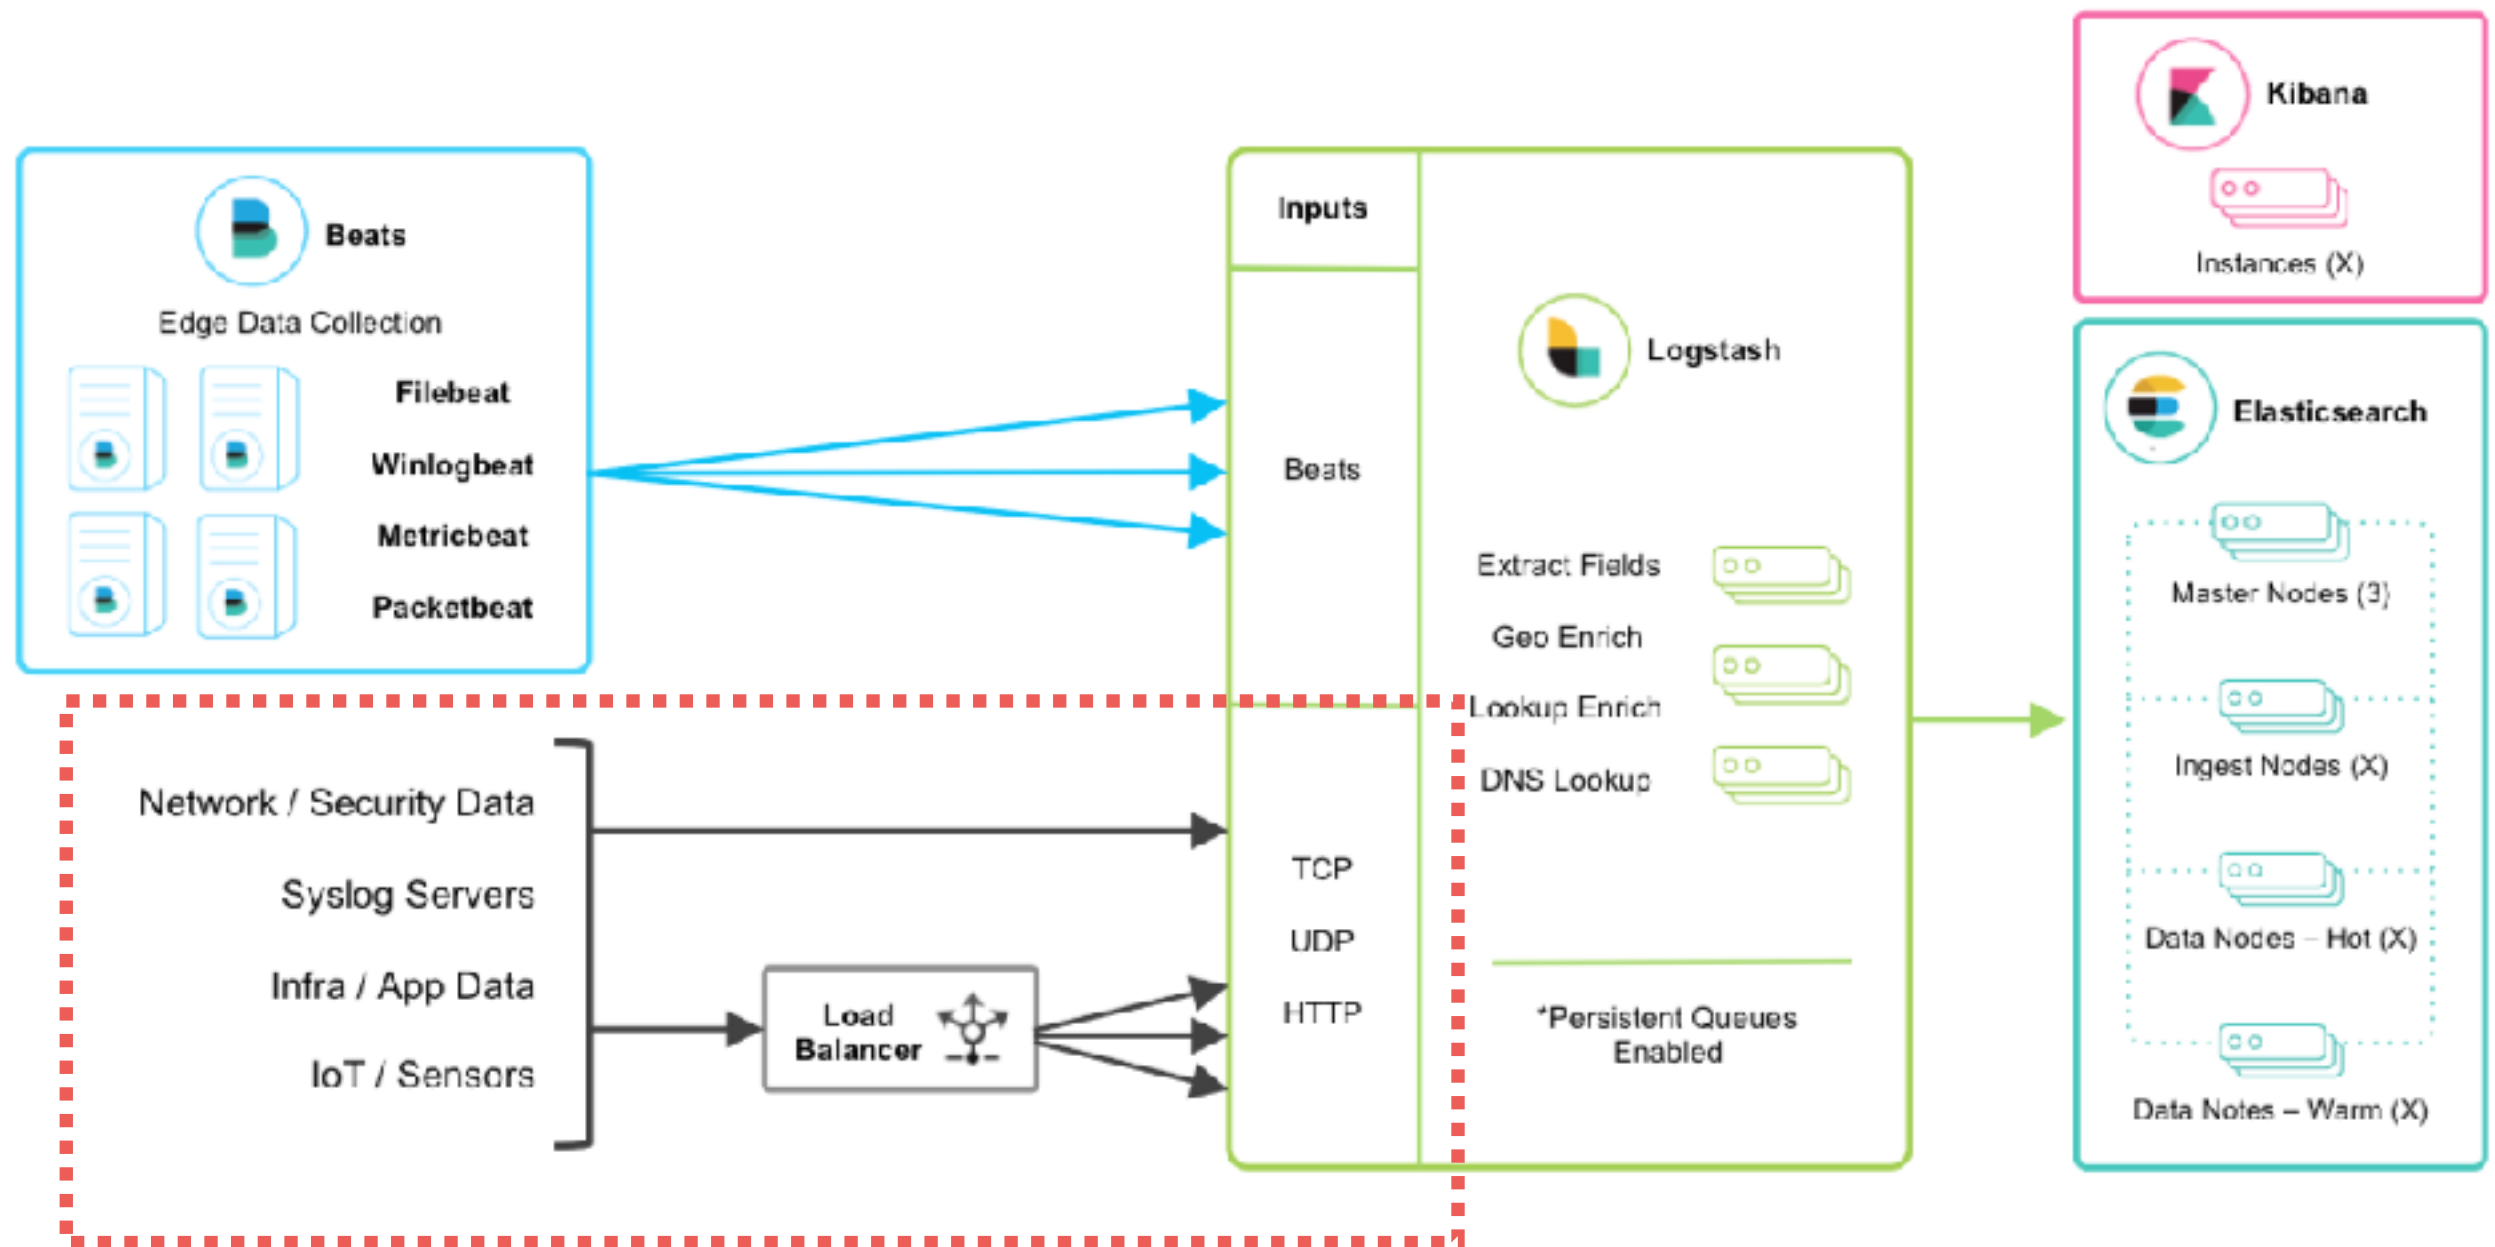
# Scaling Ingest (Add more nodes)



<https://www.elastic.co/docs/reference/logstash/deploying-scaling-logstash>



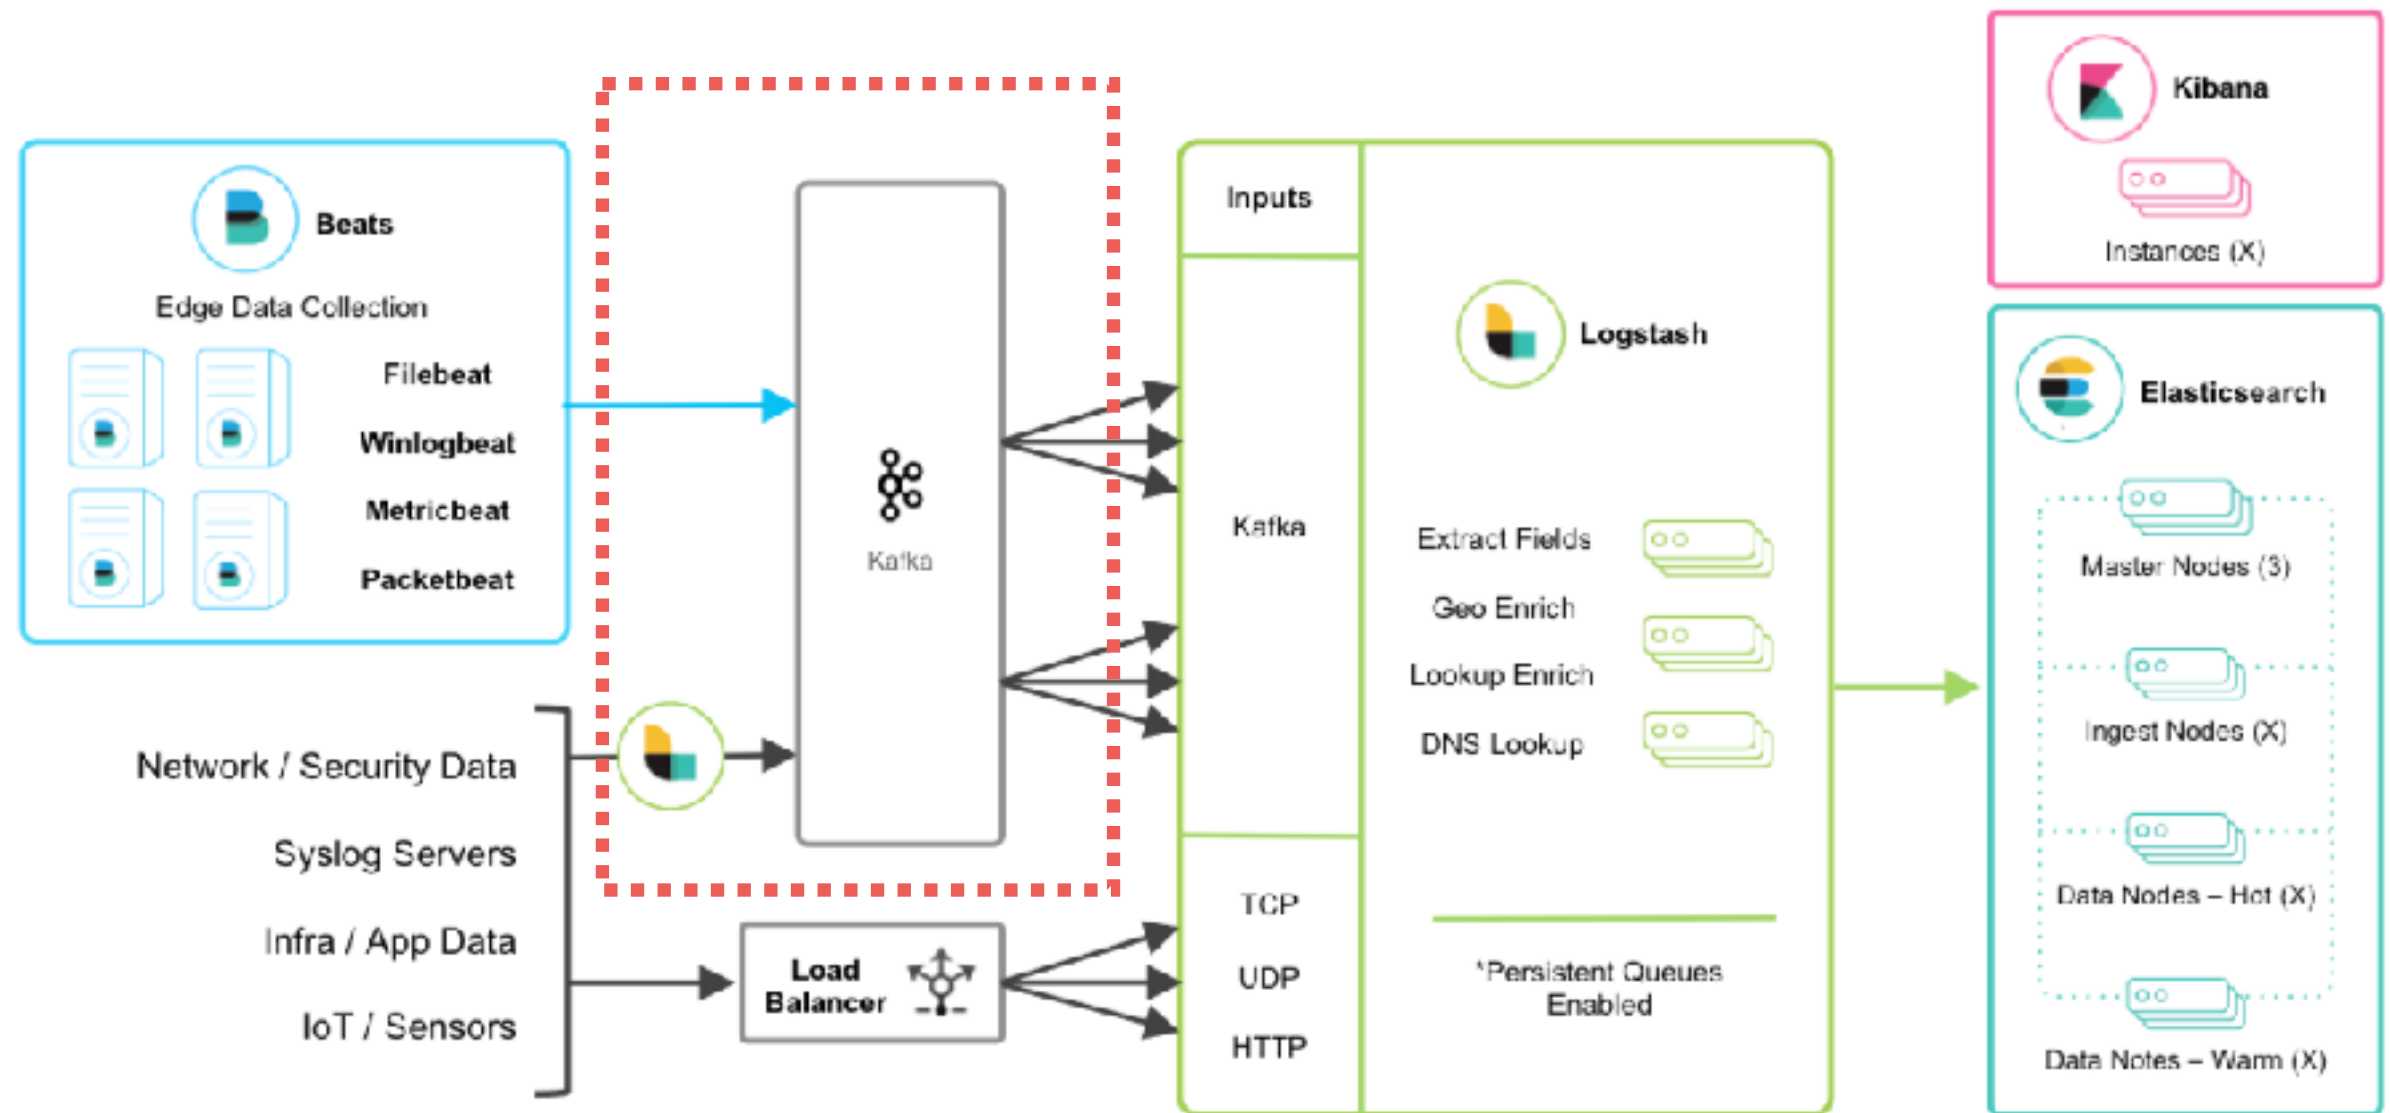
# Integrate with More data source



<https://www.elastic.co/docs/reference/logstash/deploying-scaling-logstash>



# Integrate with Messaging



<https://www.elastic.co/docs/reference/logstash/deploying-scaling-logstash>





# Messaging !!





Cloud Pub/Sub



# Why need Messaging ?

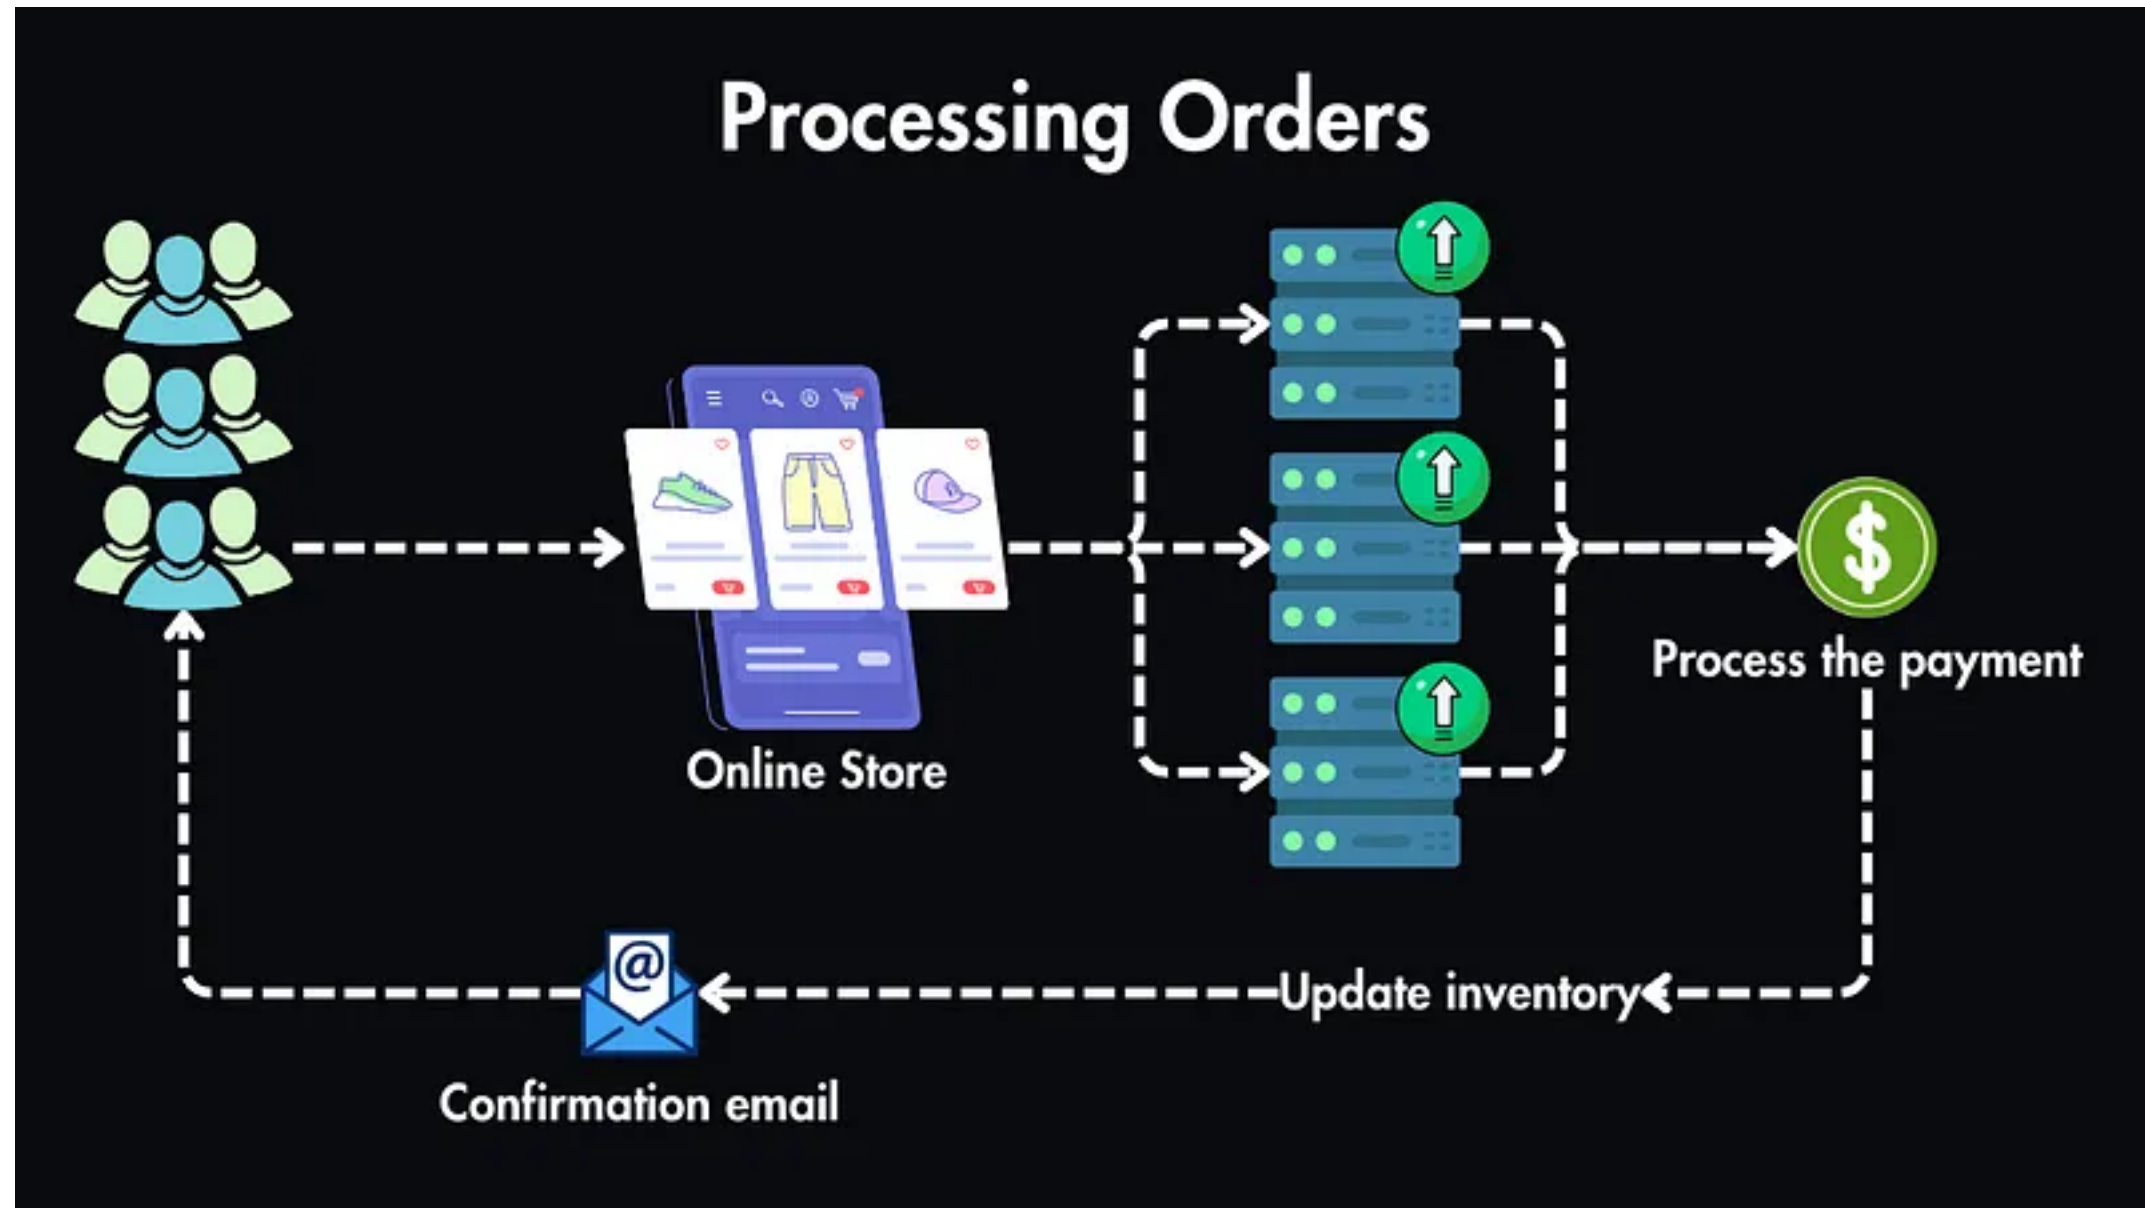
Decoupling between systems

Buffering

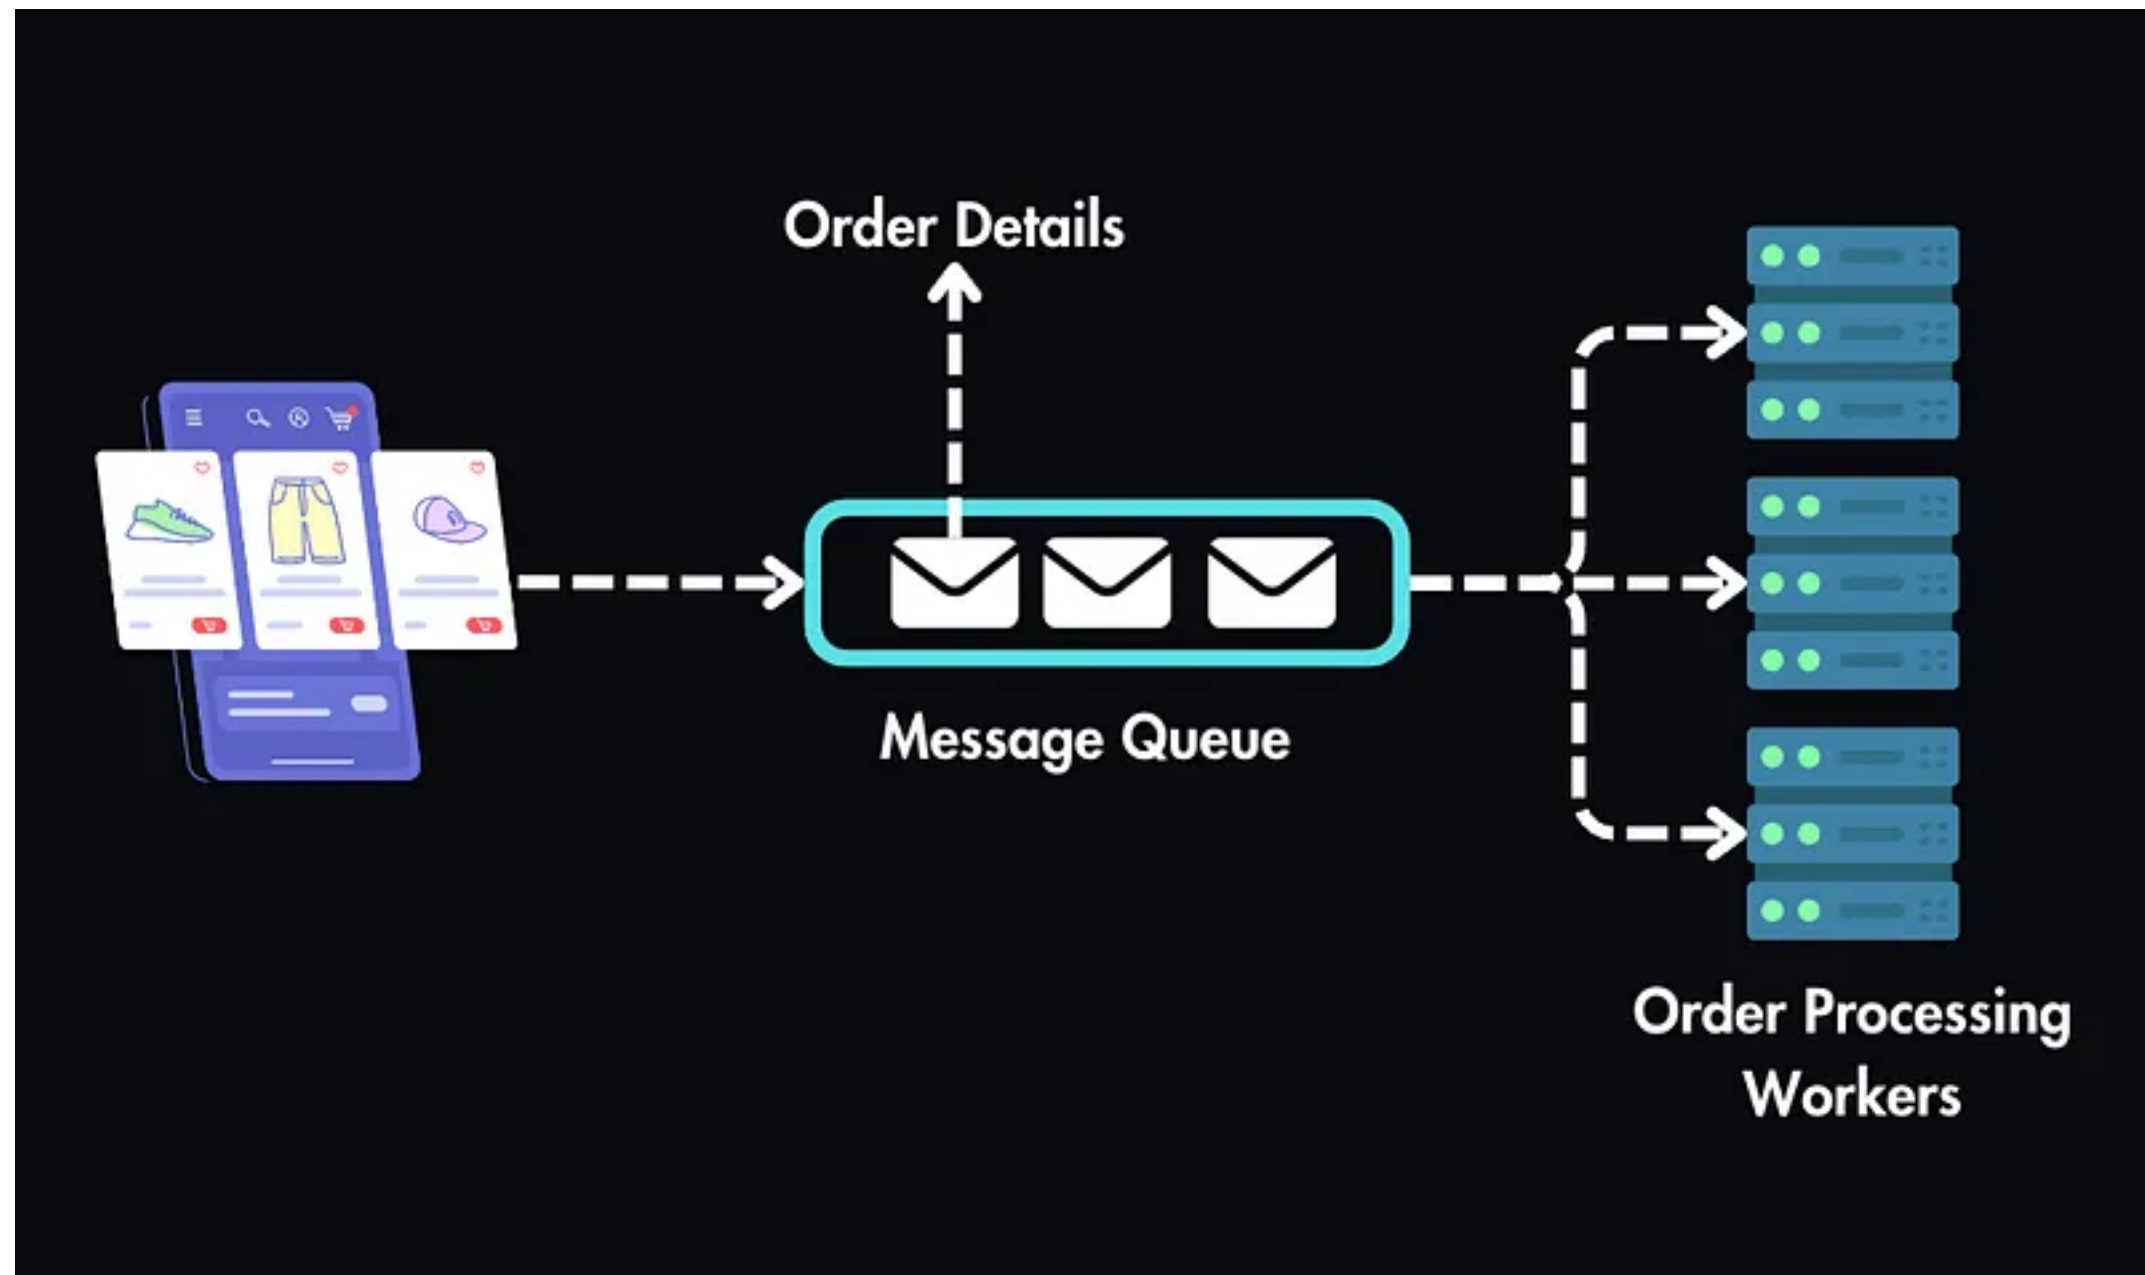
Back pressure problem



# Processing Orders



# Order processing with messaging



# Durable

If the queue crashed, that data will not lost

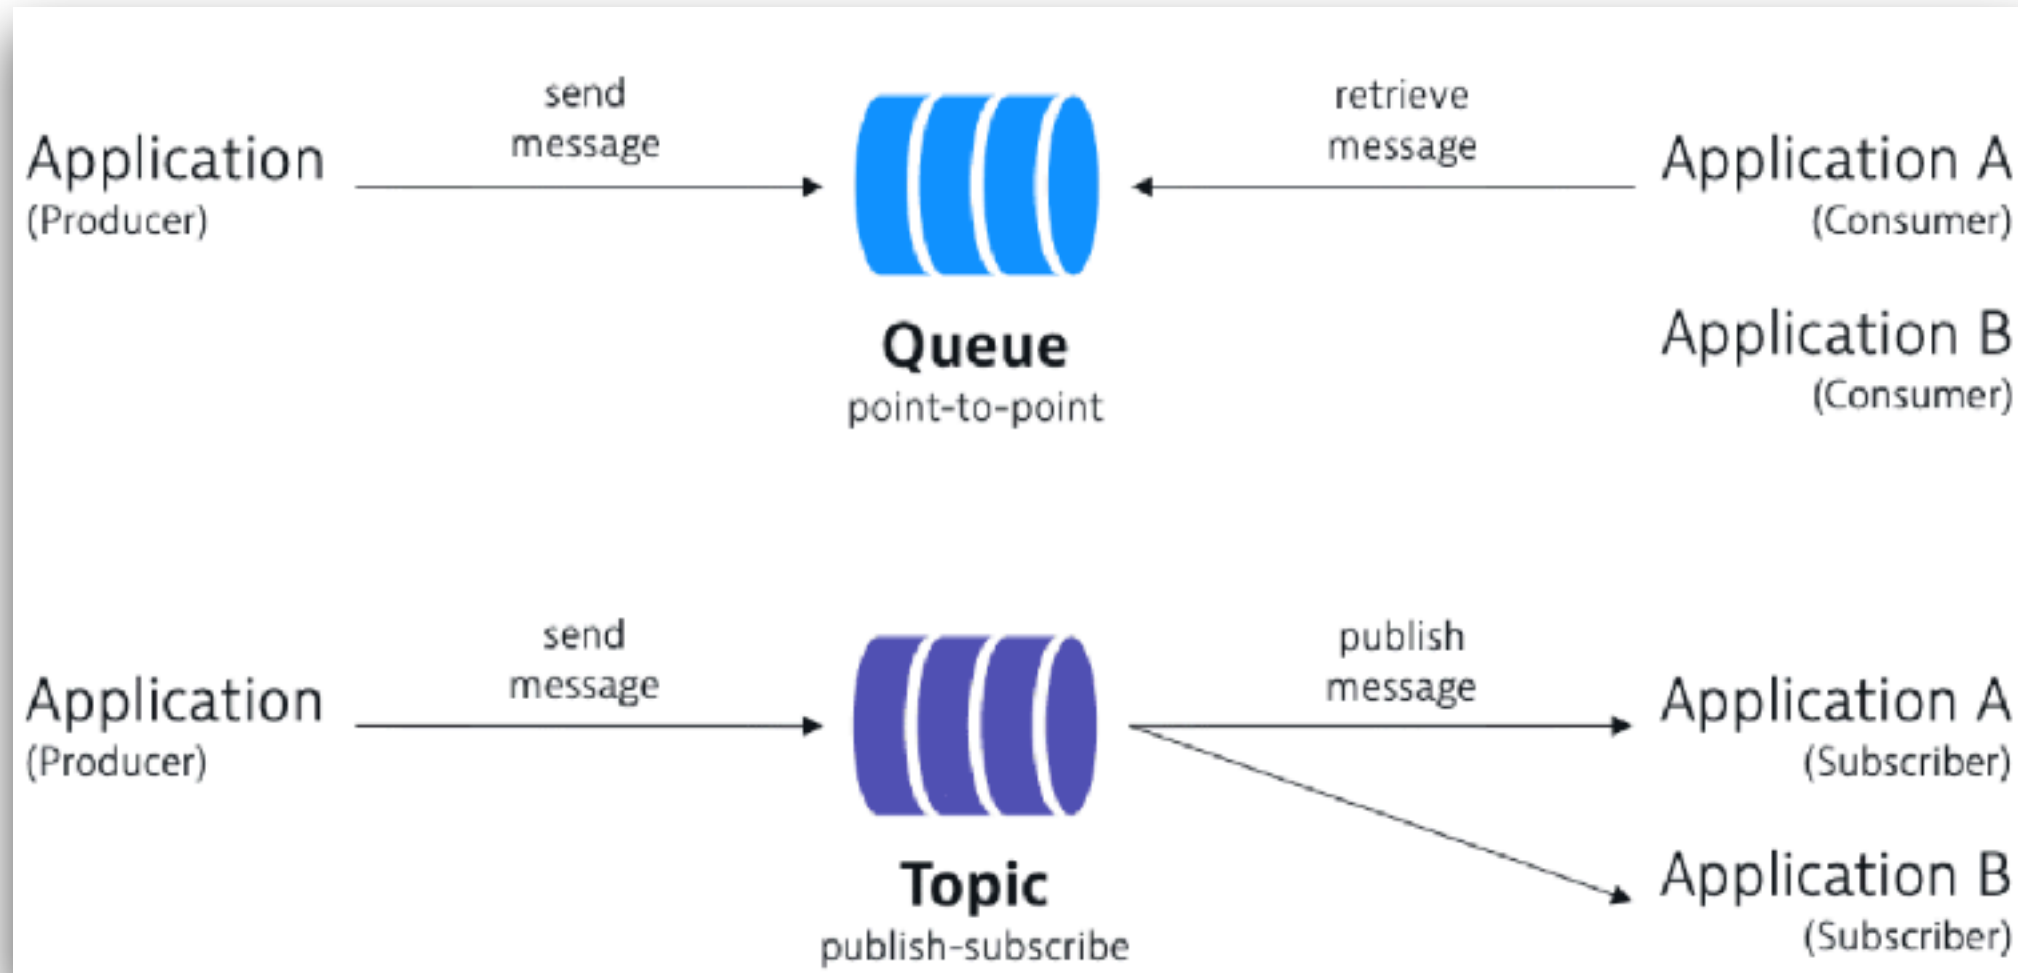


# Reliable

If consumer/worker failed, message still in queue

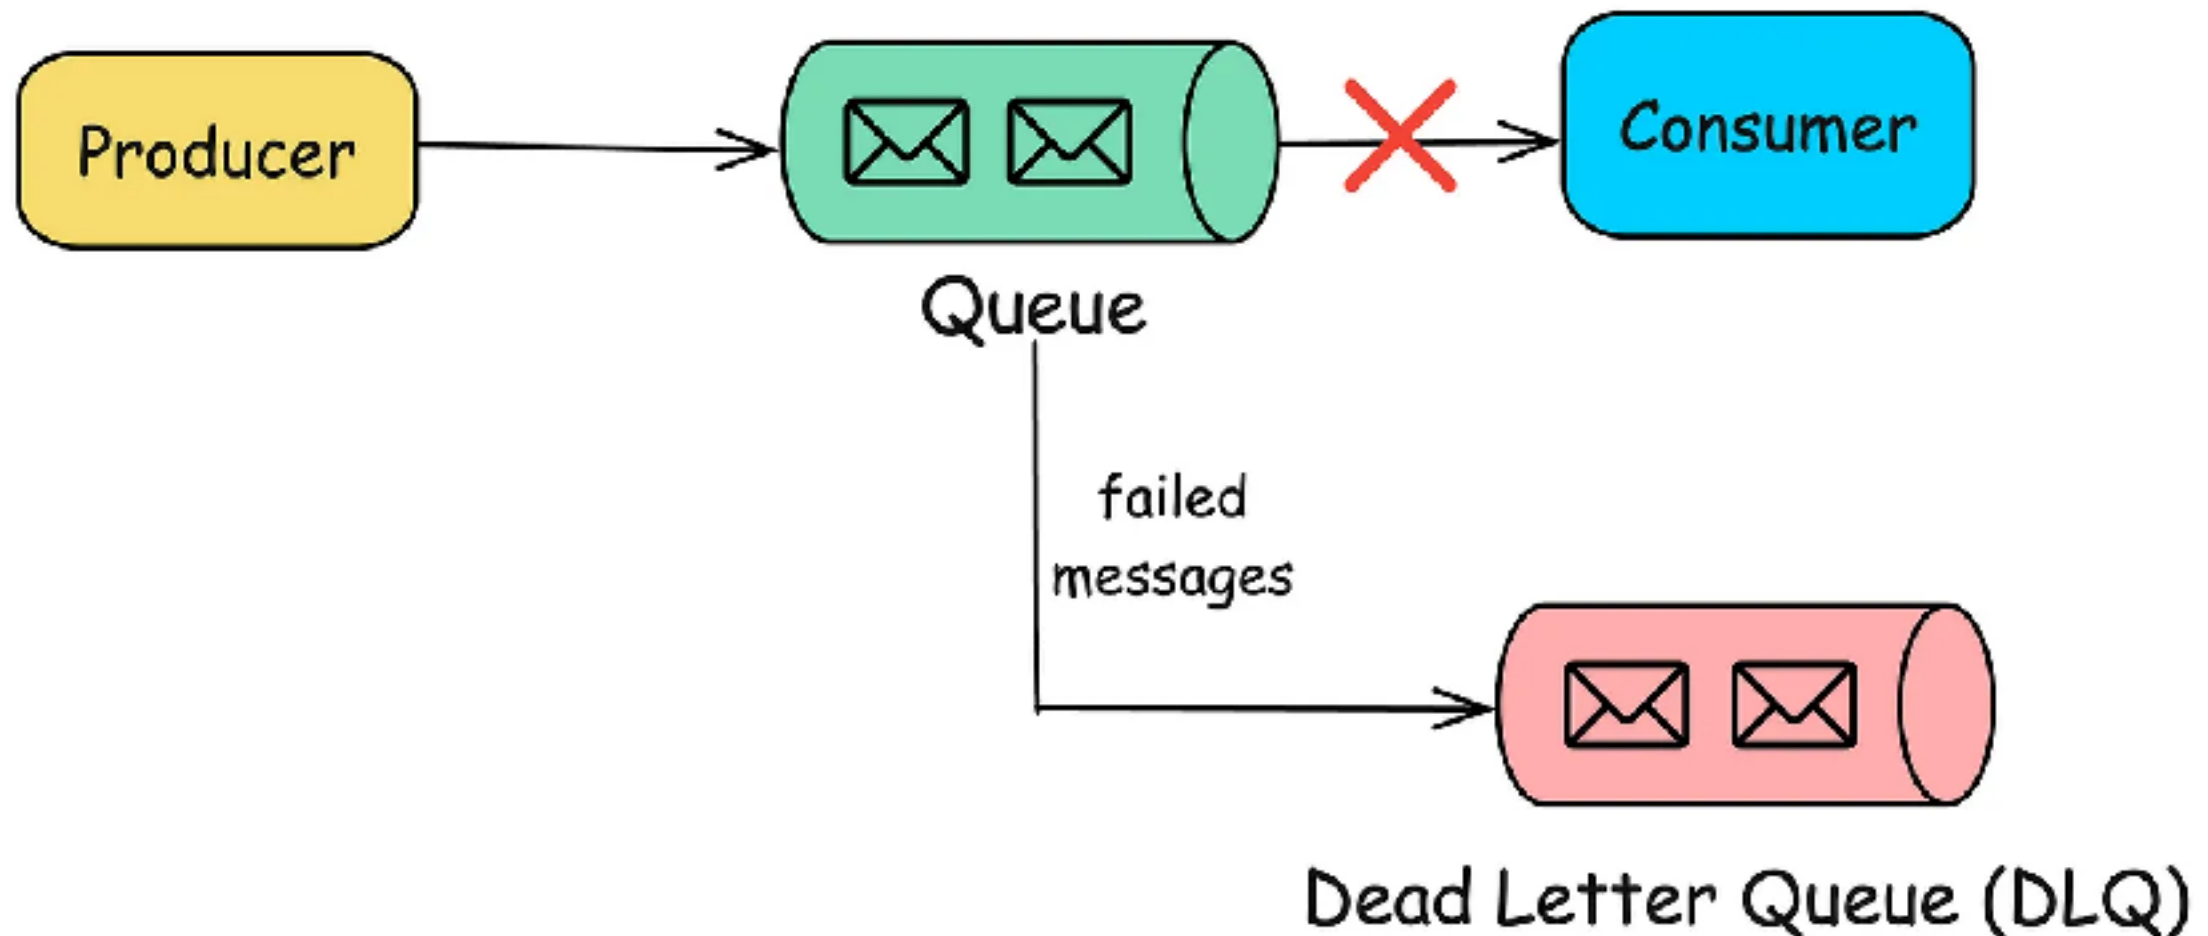


# Queue vs Topic





# Dead Letter Queue



<https://blog.algomaster.io/p/message-queues>



# Apache Kafka

<https://kafka.apache.org/>



# Kafka Use Cases

Realtime data streaming

Log aggregation

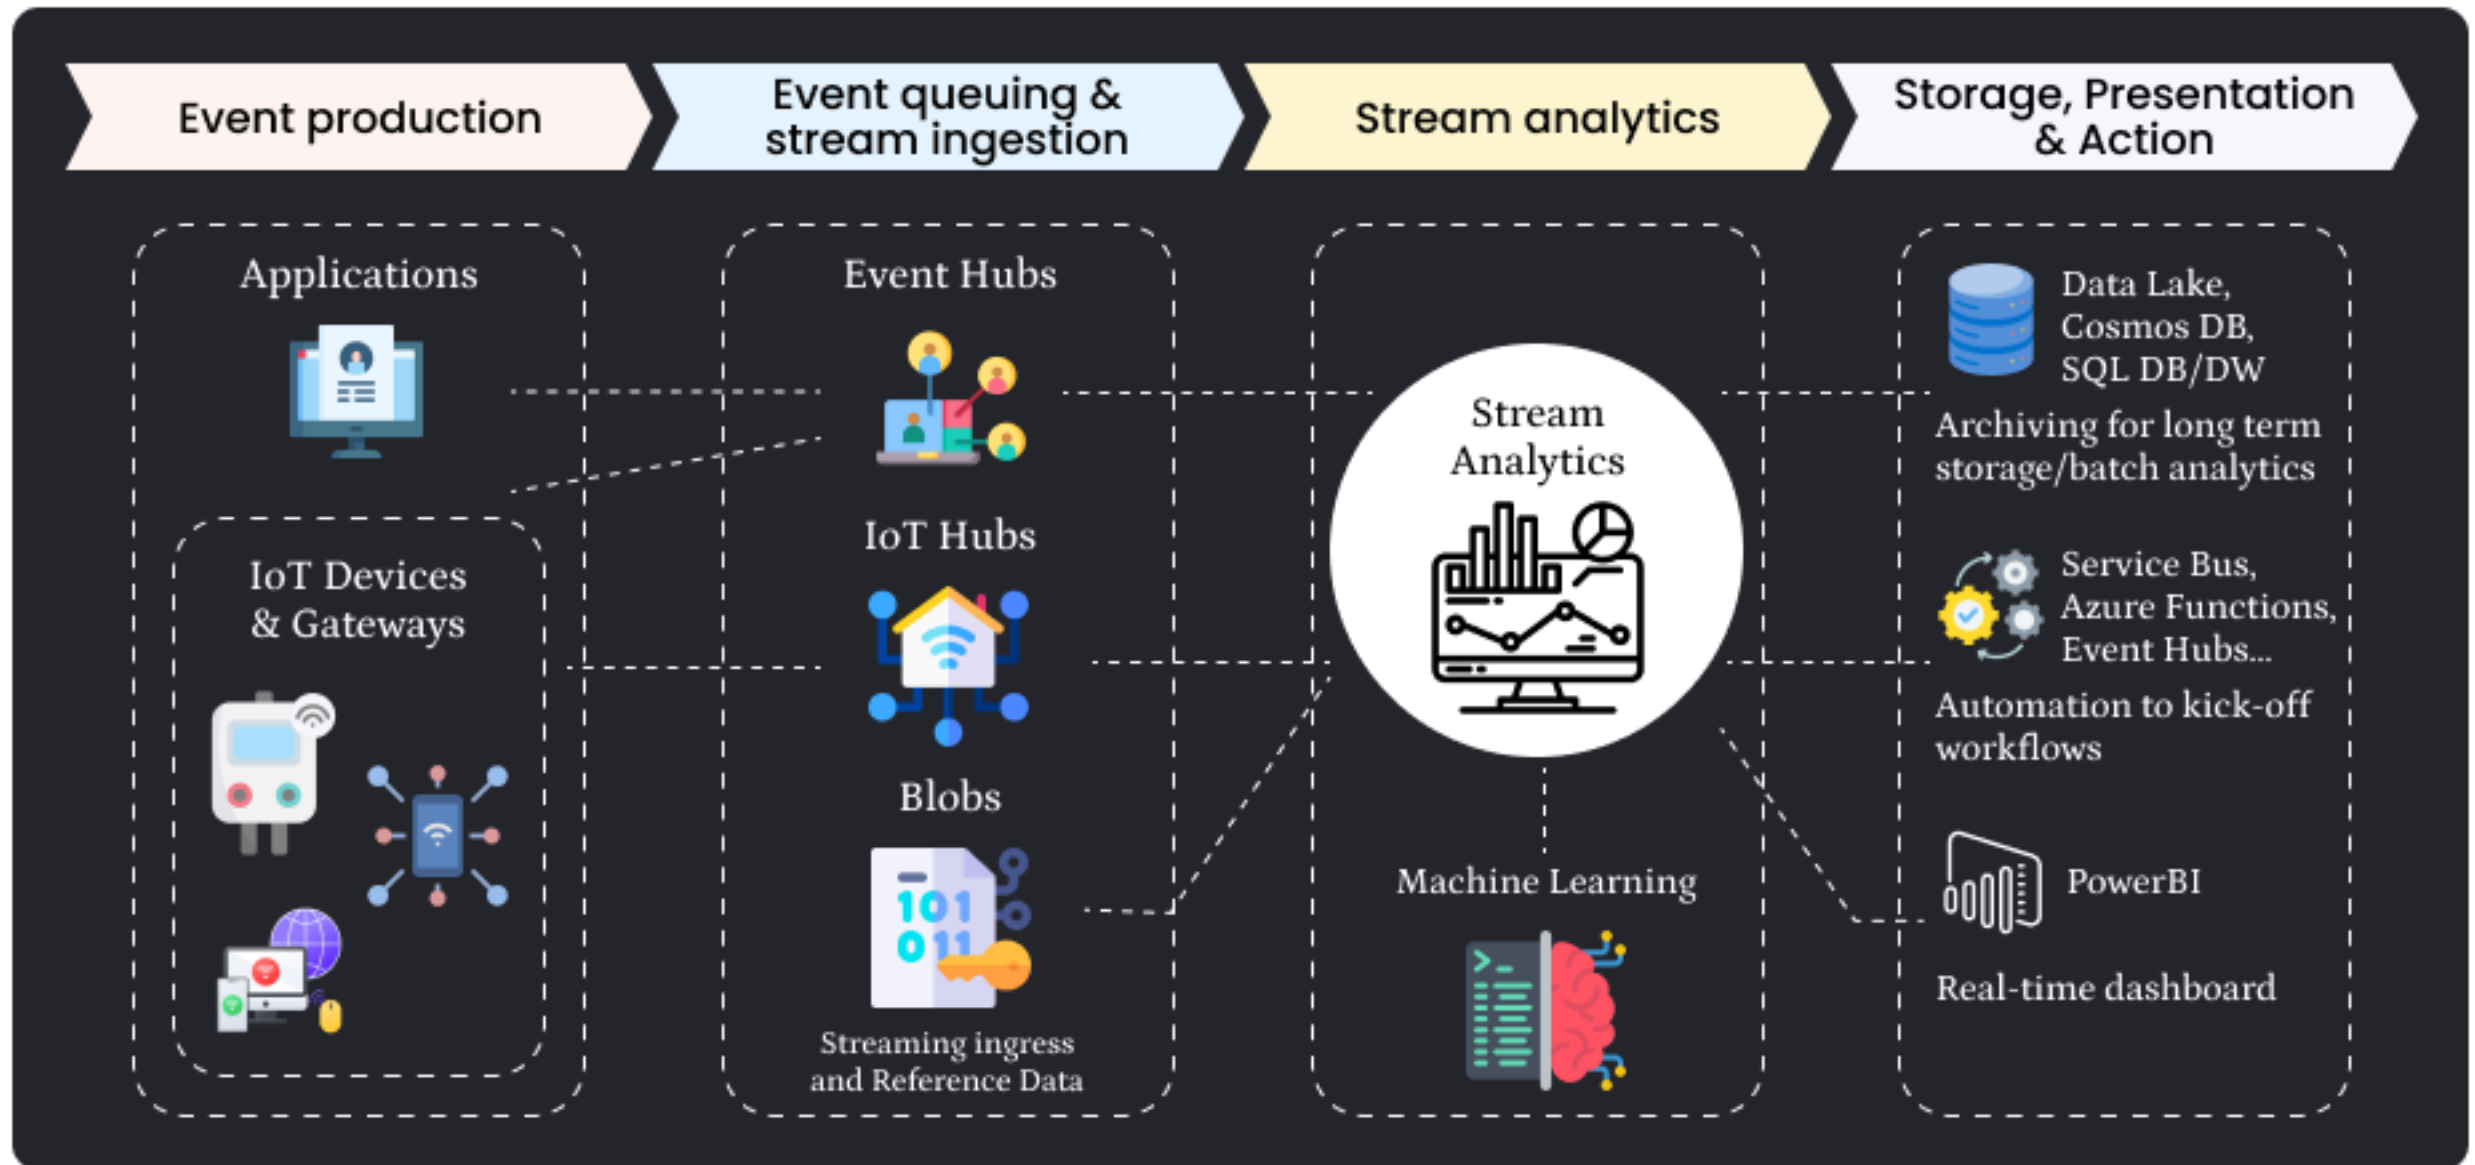
Metrics collection

Event sourcing

Streaming processing



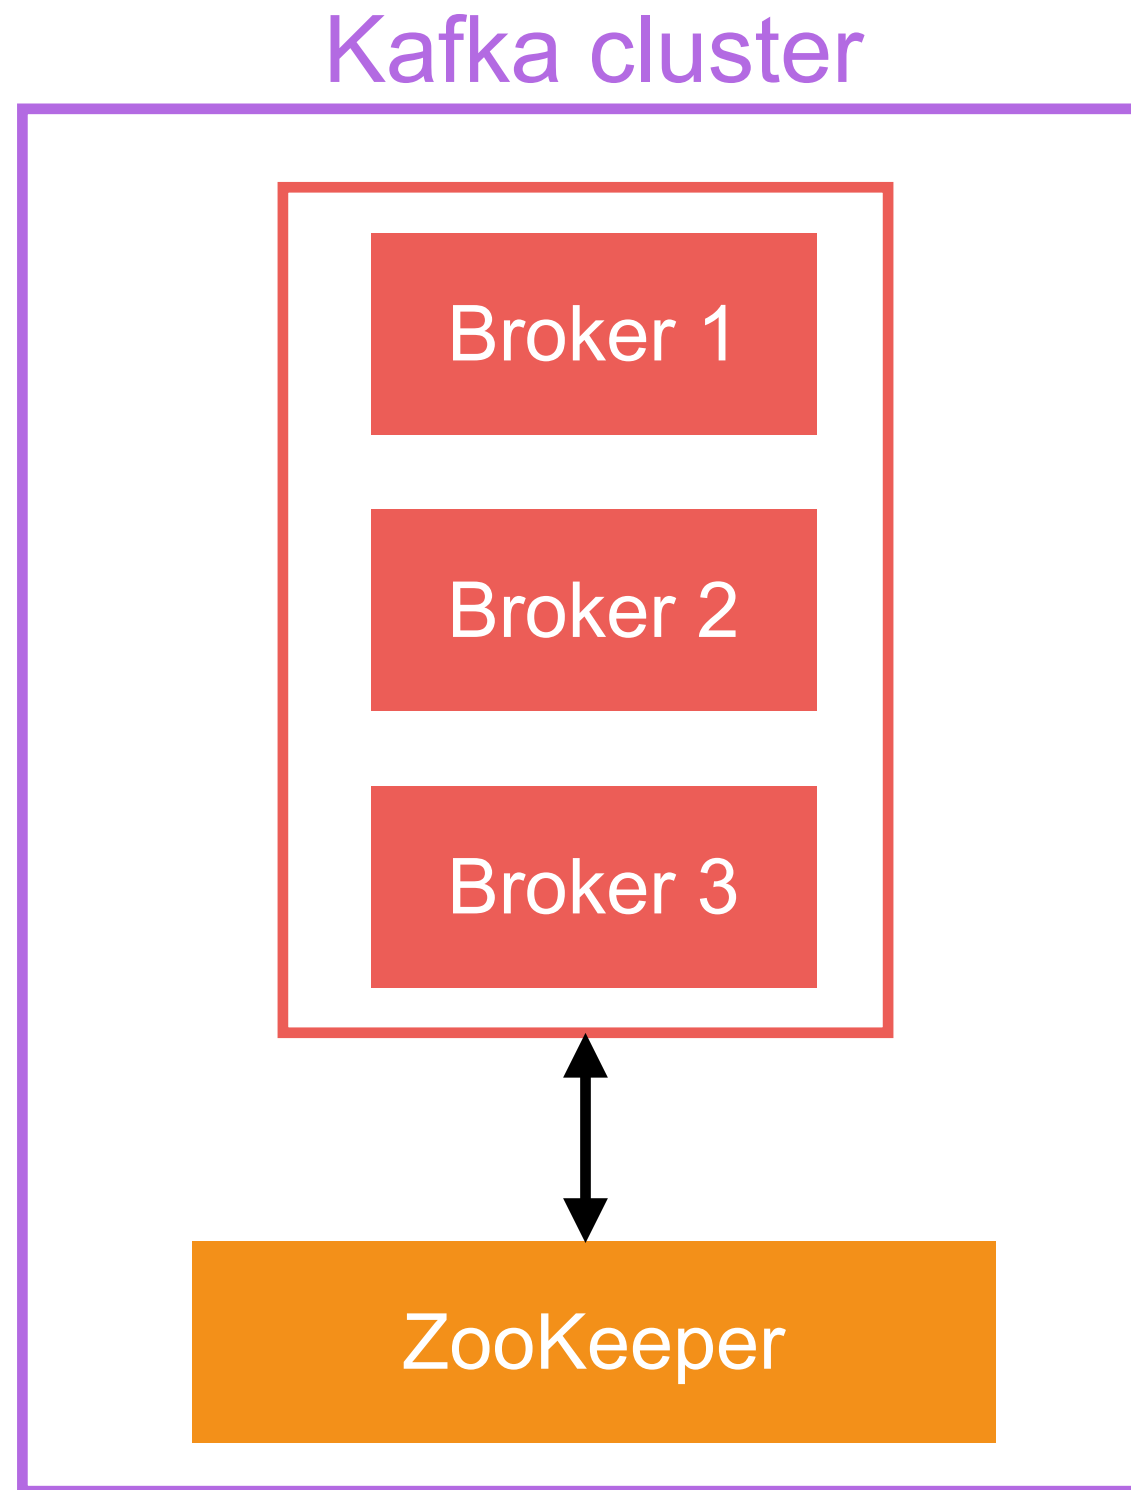
# Kafka Use Cases



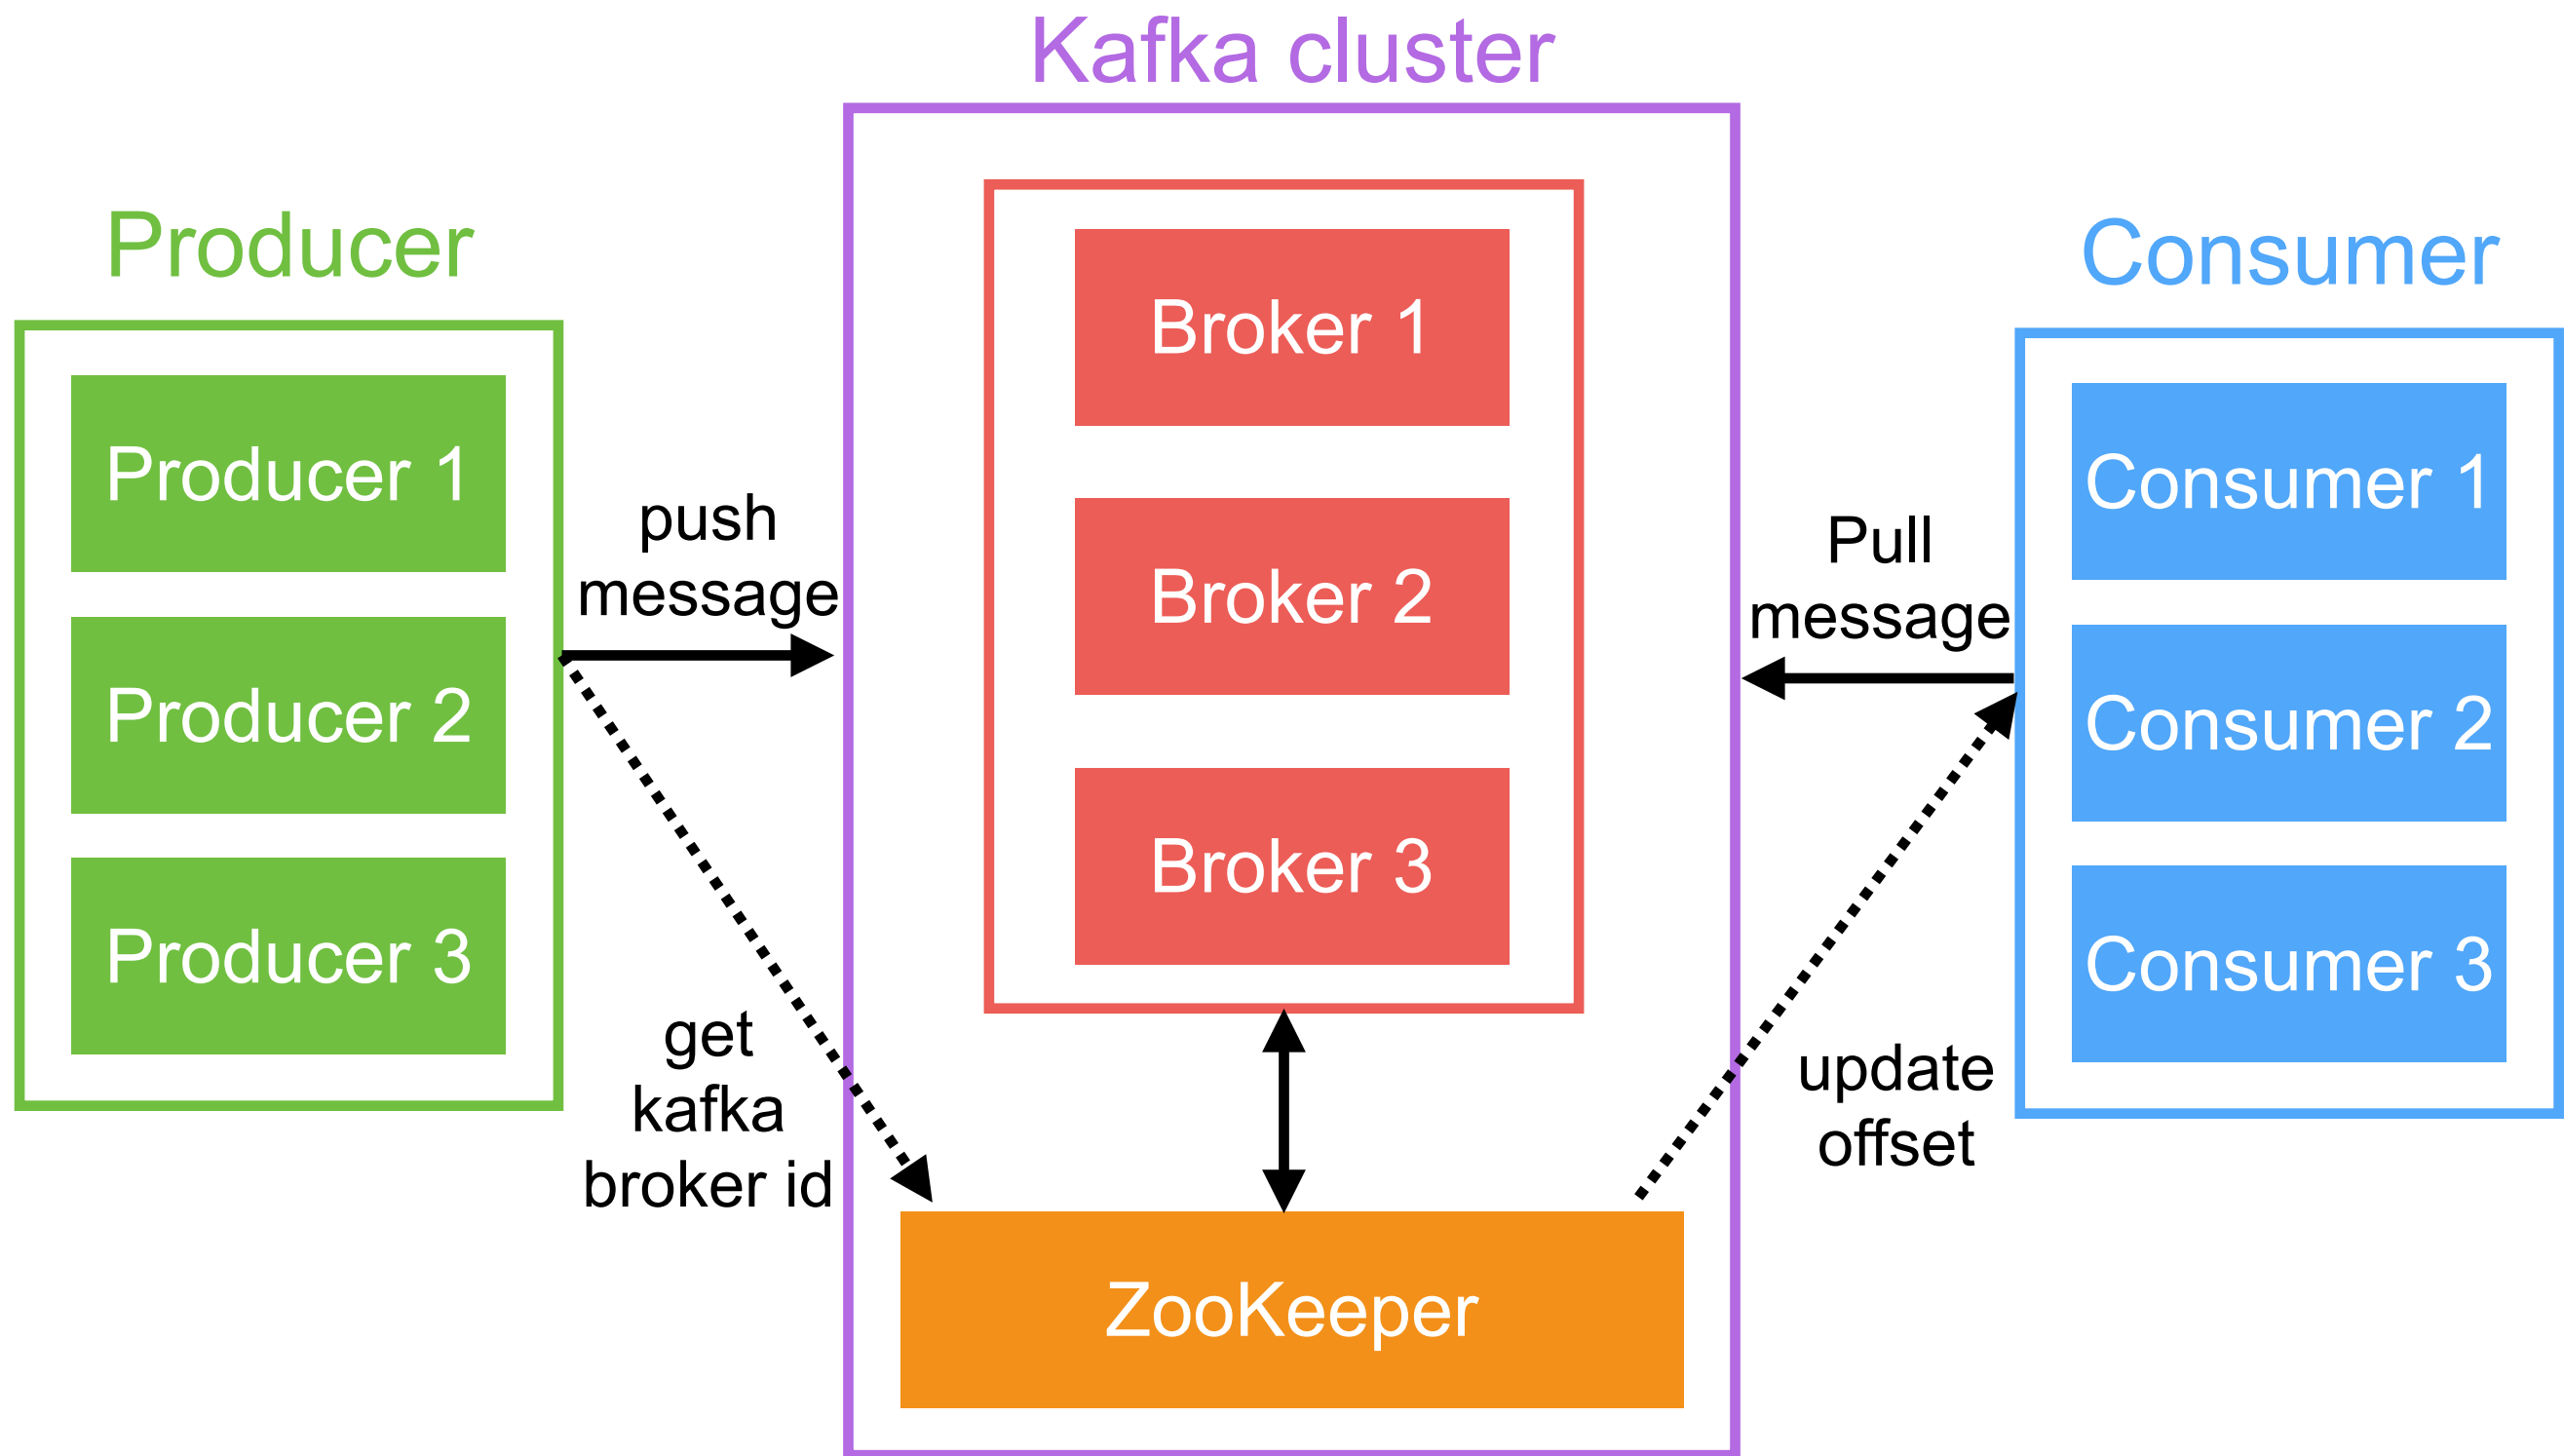
<https://www.softqubes.com/blog/apache-kafka-a-comprehensive-guide-to-real-time-data-streaming-and-processing/>



# Kafka Architecture



# Kafka Architecture



# APACHE KAFKA

# 4.0

2025/04/18



# Kafka 4.0

Remove Apache Zookeeper, use Kraft



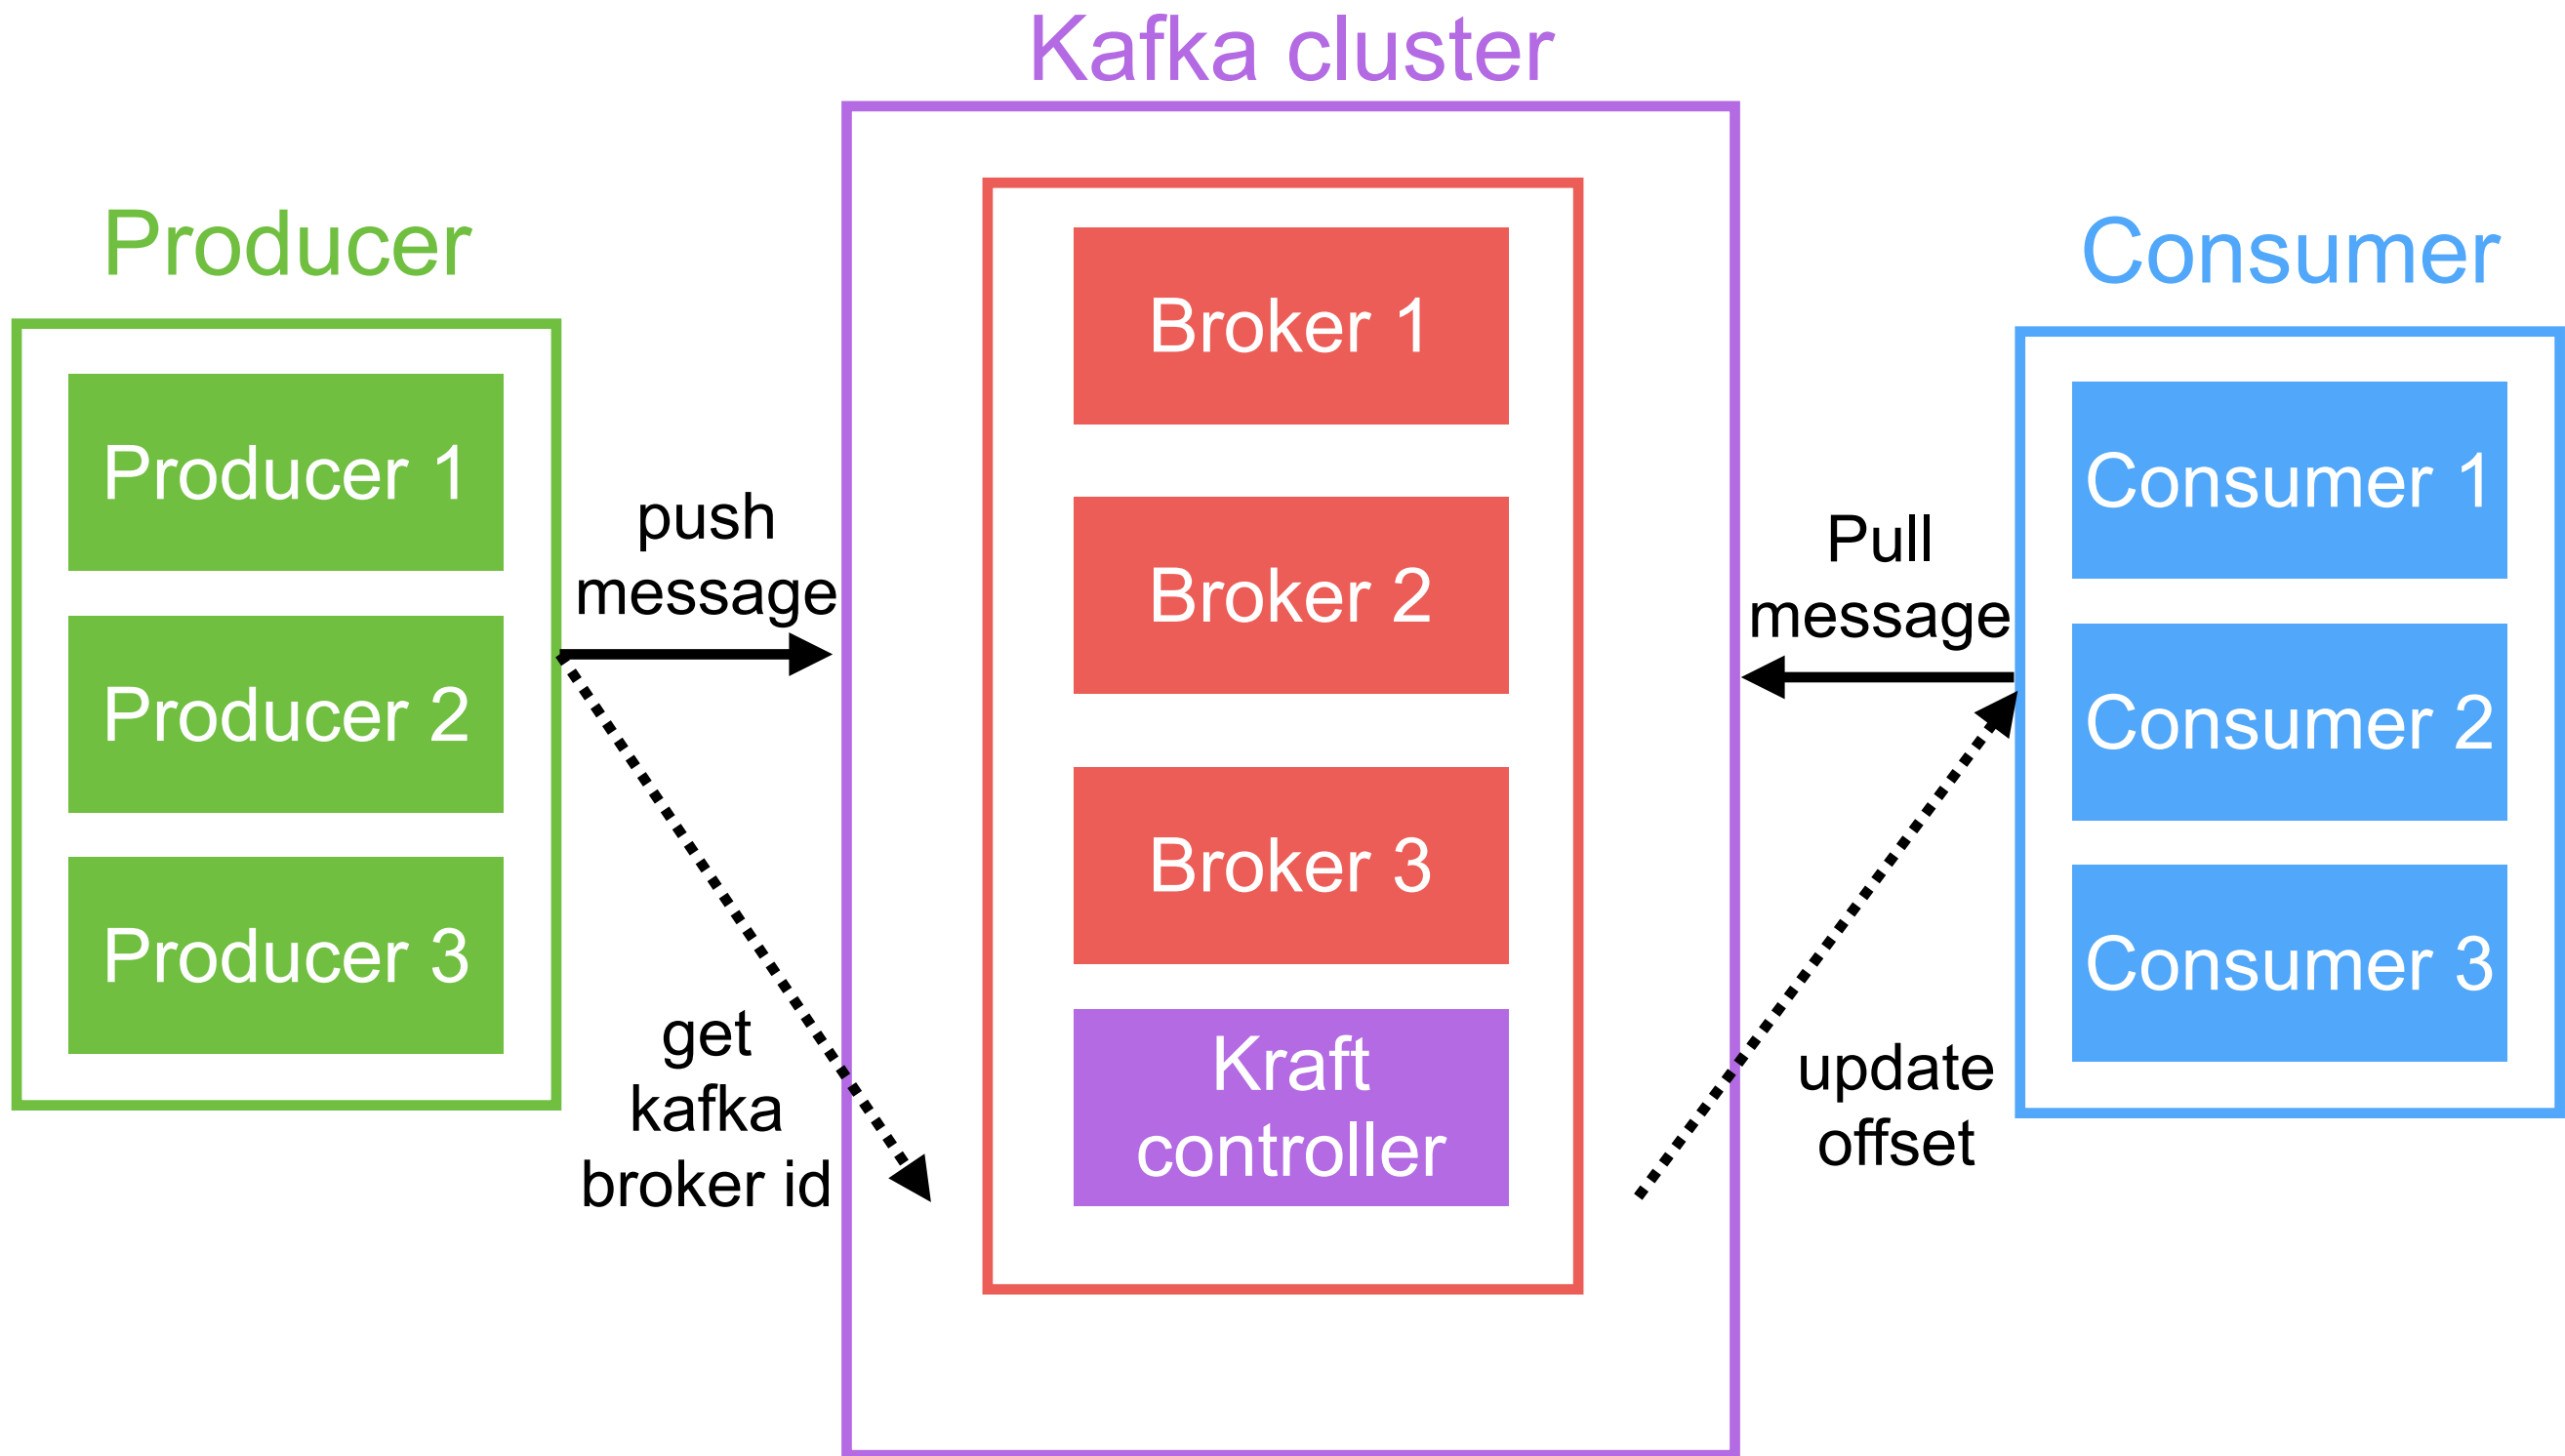
APACHE  
**ZooKeeper**<sup>TM</sup>

<https://zookeeper.apache.org/>

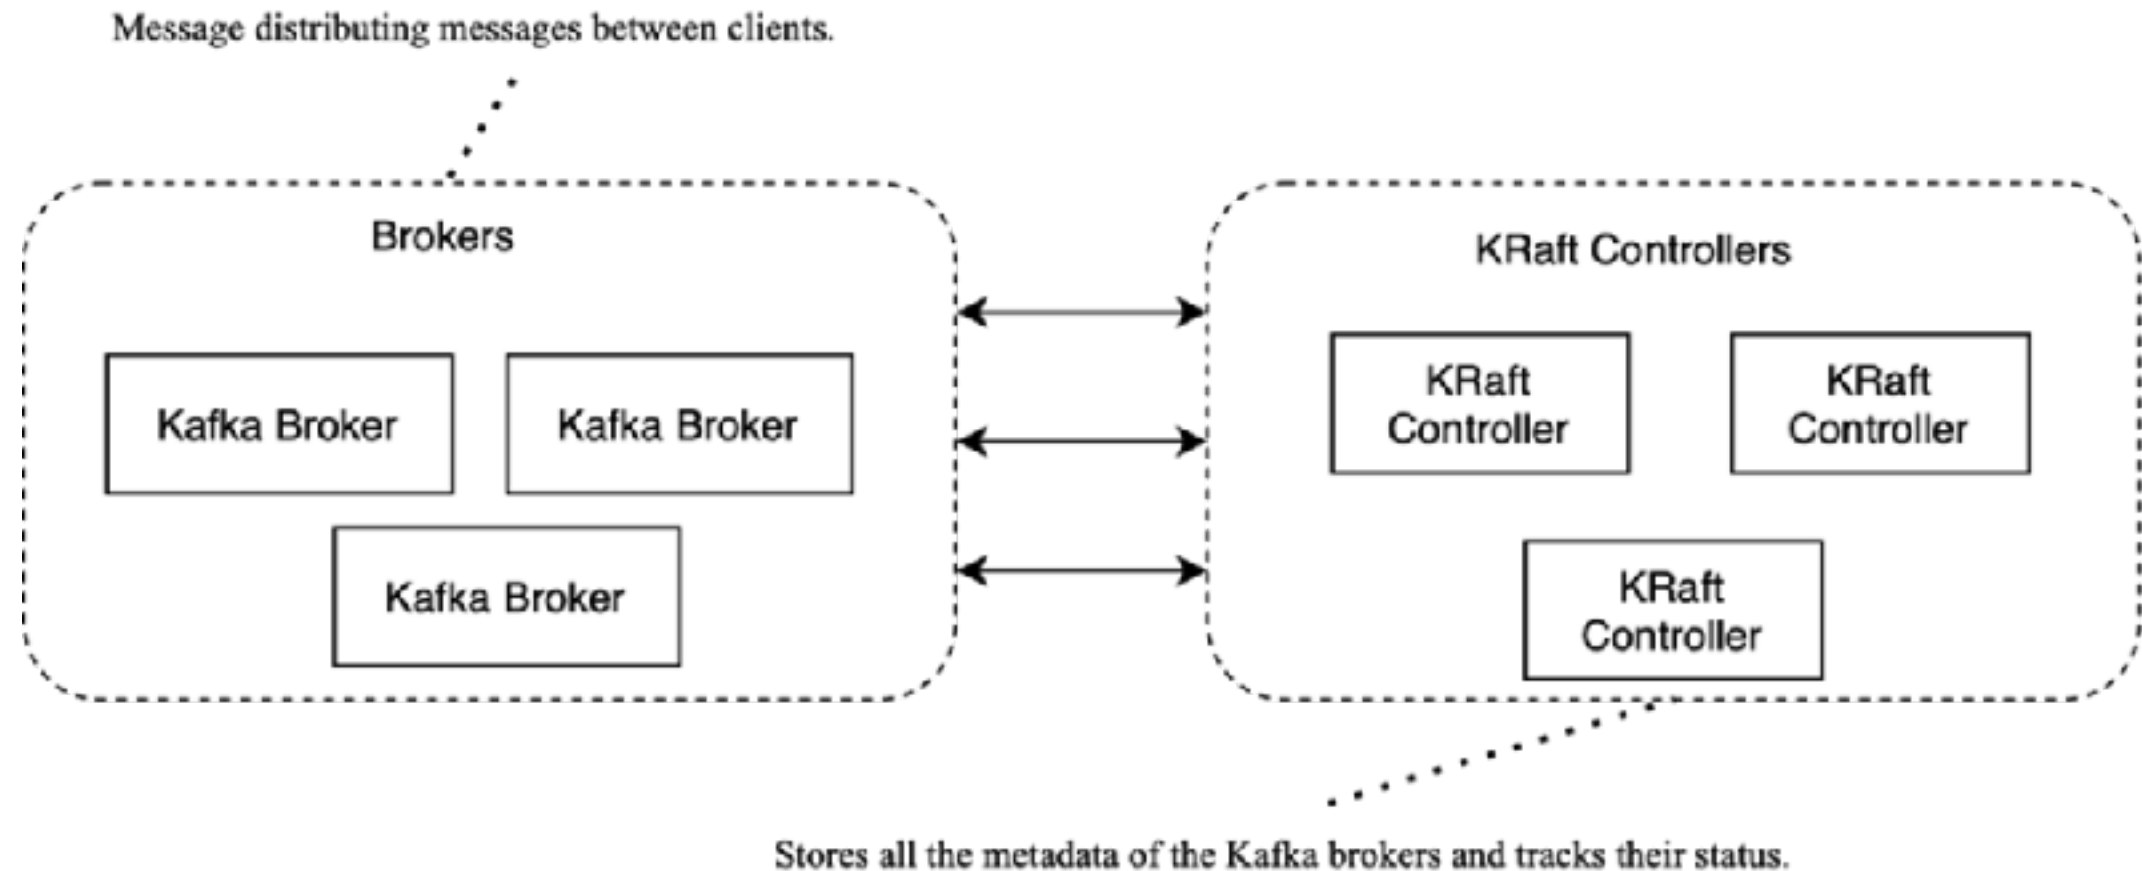




# Kafka Architecture 4.0



# Kafka broker and KRaft controller



# Kafka Fundamentals

Event/record/message

Producers

Consumers

Topics

Partitions

Broker

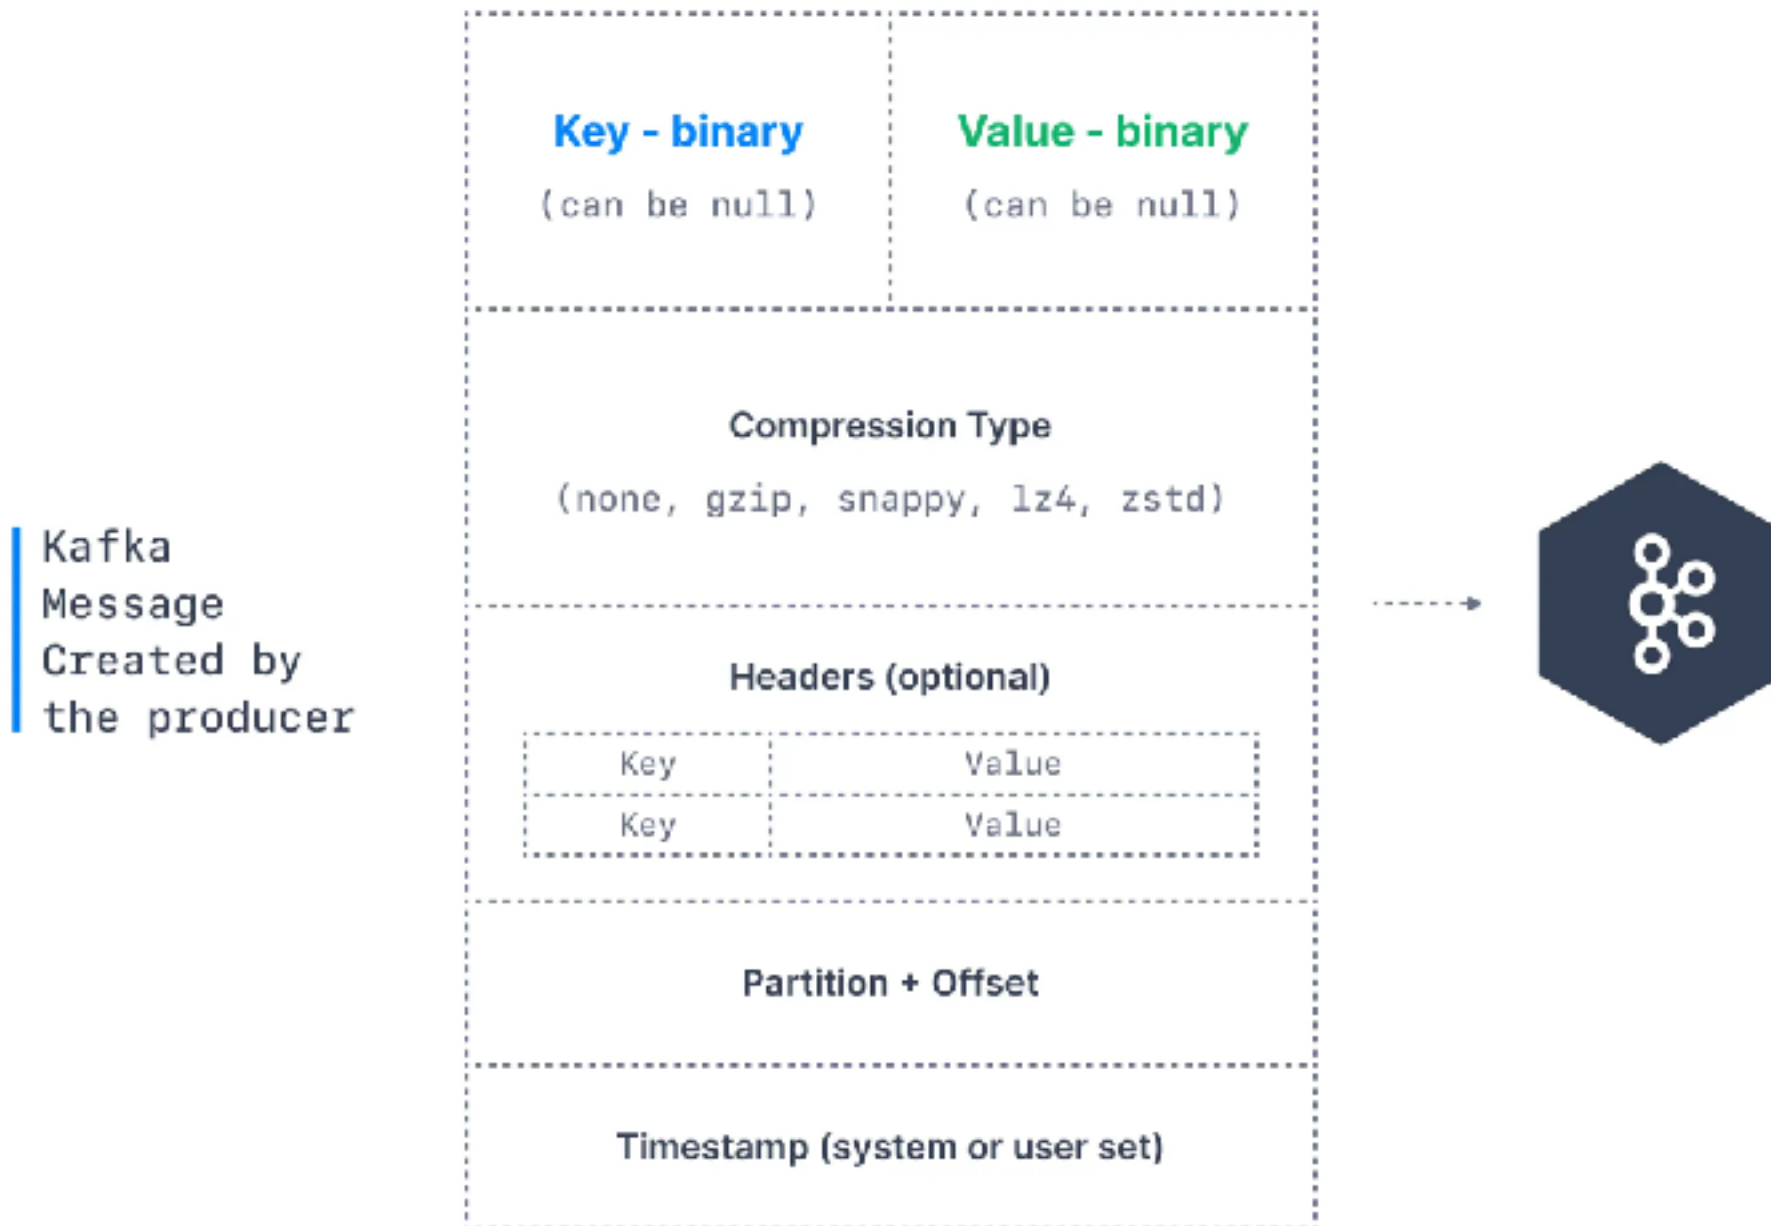
Consumer groups

Cluster/Replicate

Offset (read and write)



# Event Anatomy



# Topics, partitions and offsets

## Topics

Stream of data

Similar to a table in database

You can have many topics as you want

A topic is identified by **name**




# Topics, partitions and offsets

## Partitions


Topics are split in partitions

Each partition is **ordered**

Each message in partition get incremental id (**offset**)



# Topics, partitions and offsets



Order is guaranteed only within partition

Data is keep in limited time (**Default = 1 week**)

Data is **immutable** (can't be changed)

Data is assigned **randomly** to a partition



# More partitions ?

Higher is your data throughput

More open files

Longer downtime

More RAM is consumed by clients

4000  
per broker

200,000  
per cluster

50 brokers  
per cluster

<https://api7.ai/blog/why-kafka-needs-an-api-gateway>





# Brokers and topics

## Brokers

Kafka cluster is composed of multiple brokers (servers)

Each broker is identified by ID (integer)

Each broker contains certain topic partitions

After connecting any broker (**bootstrap broker**),  
you will be connected to the entire cluster

Broker 1

Broker 2

Broker 3



# Kafka broker discovery

Every Kafka broker is called “**bootstrap server**”

You only need to connect to one broker, and you will be connected to the entire cluster

Broker 1  
(bootstrap)

Broker 2  
(bootstrap)

**Kafka cluster**

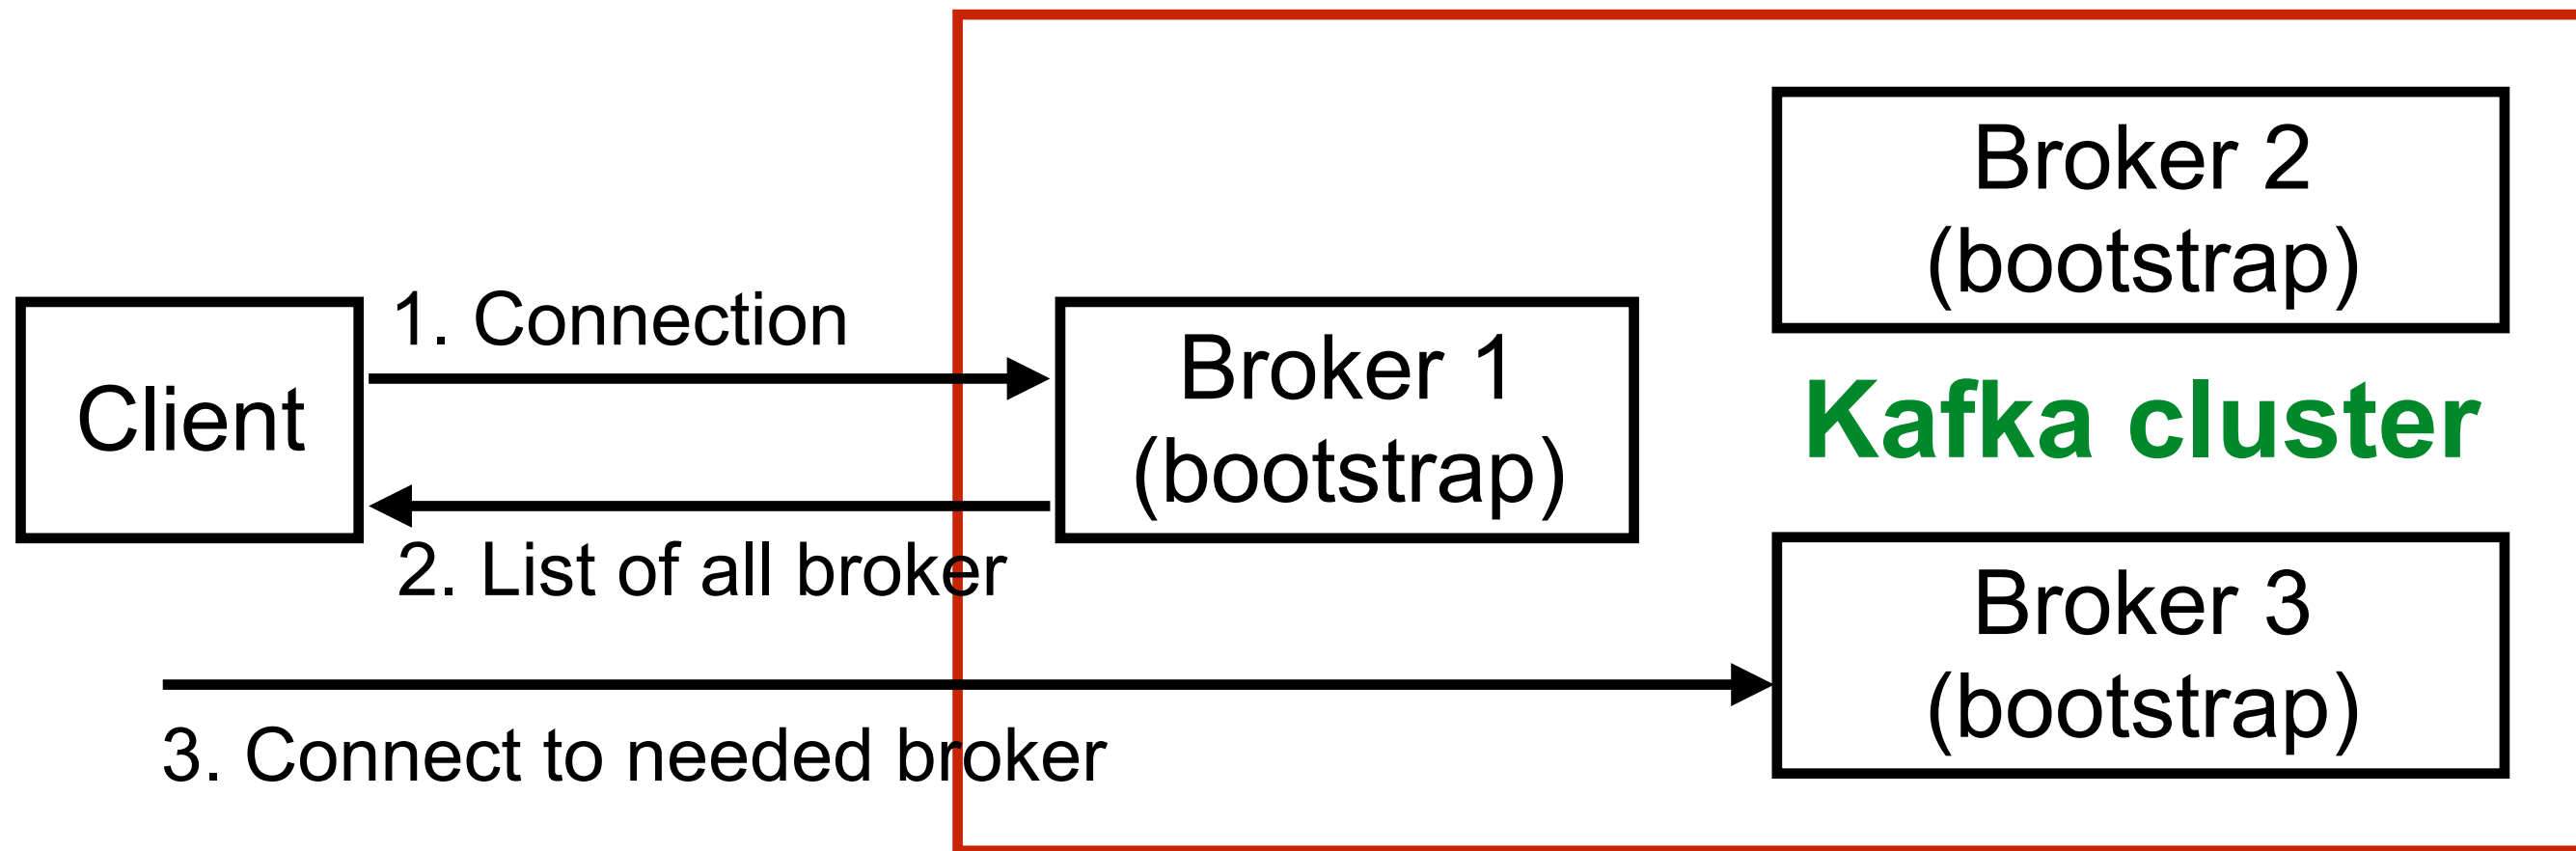
Broker 3  
(bootstrap)

Broker 4  
(bootstrap)



# Kafka broker discovery

Each broker knows about all brokers, topics and partitions (**metadata**)



# Apache Zookeeper

Kafka < 4.0 can not work without Zookeeper !!



<https://zookeeper.apache.org/>



# Apache Zookeeper

Zookeeper **manages** brokers

Zookeeper help in performing **leading election** for partitions

Zookeeper sends **notifications** to Kafka in case of changes (new topic, broker die)

Zookeeper by design operates with a odd number of servers (1, 3, 5, 7)



# Data in Zookeeper

| Kafka's data operations | Format                                     |
|-------------------------|--|
| Broker metadata         | ID, hostname                               |
| Topic metadata          | Topic name, partition count, replica count |
| Partition assignment    | Leader partition                           |
| Consumer group metadata | Consumer group name                        |
| Cluster metadata        | Active broker, list of topics, partitions  |
| Leader election         | Select leader of partition                 |




# Zookeeper architecture

Leader for write  
Follows for read



# Zookeeper architecture





# Brokers and topics

## Brokers

Good number to start id **3 brokers**

Broker 1

Broker 2


Broker 3



# Brokers and topics

## Example


Topic A with 3 partitions



# Brokers and topics

## Example

Topic B with 2 partitions



# Topic with replication factor

Topic should have a replica factor  $> 1$  (2-3)

**replica factor  $<$  no. of brokers**

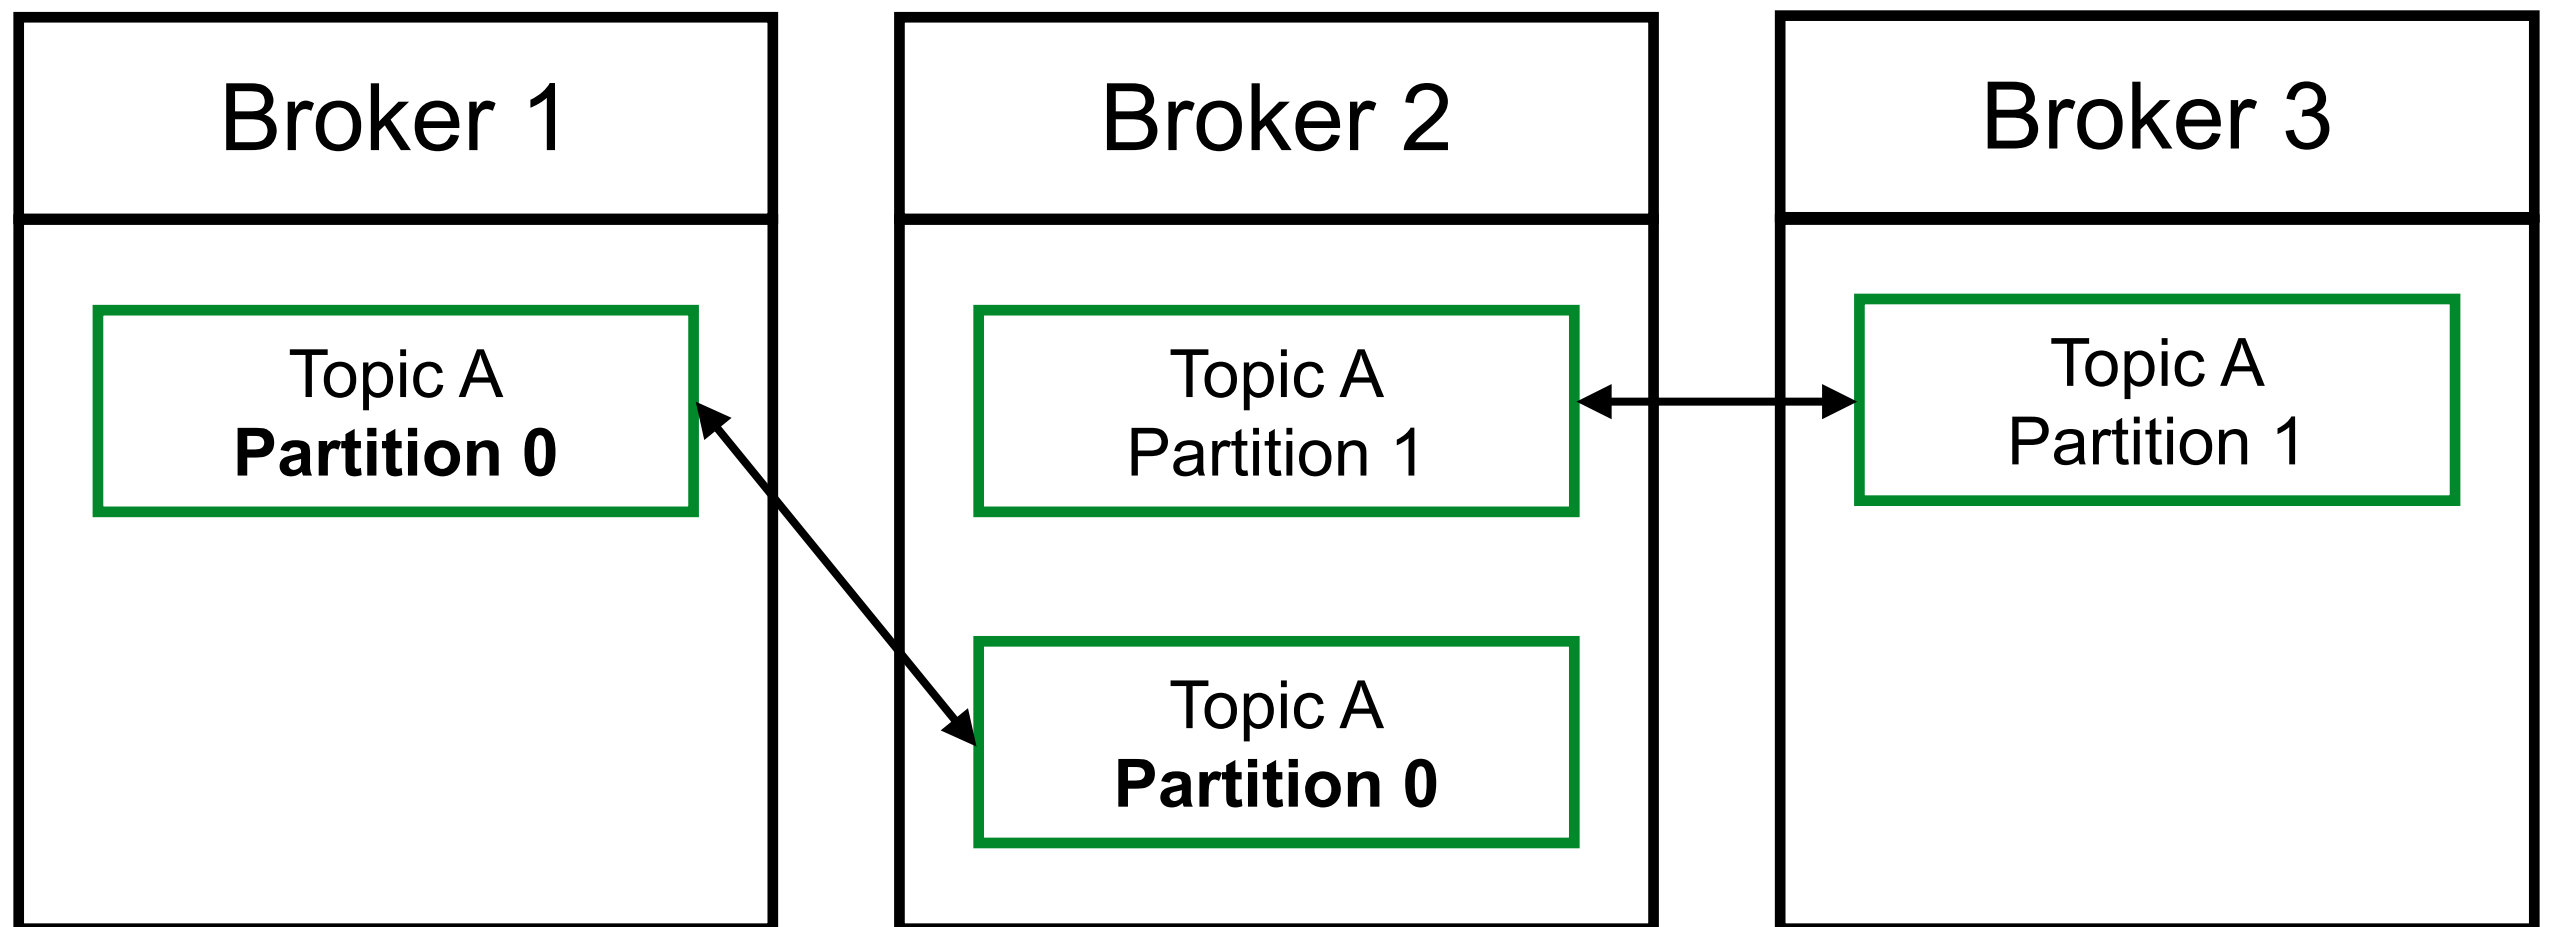
This way if a broker down, another broker can serve the data



# Topic with replication factor

## Example

Topic A with 2 partitions and replication factor = 2



# Leader for a partition

At any time only **one Broker** can be leader for partition

**Only leader partition** can receive and serve data for a partition

Other brokers will synchronize the data

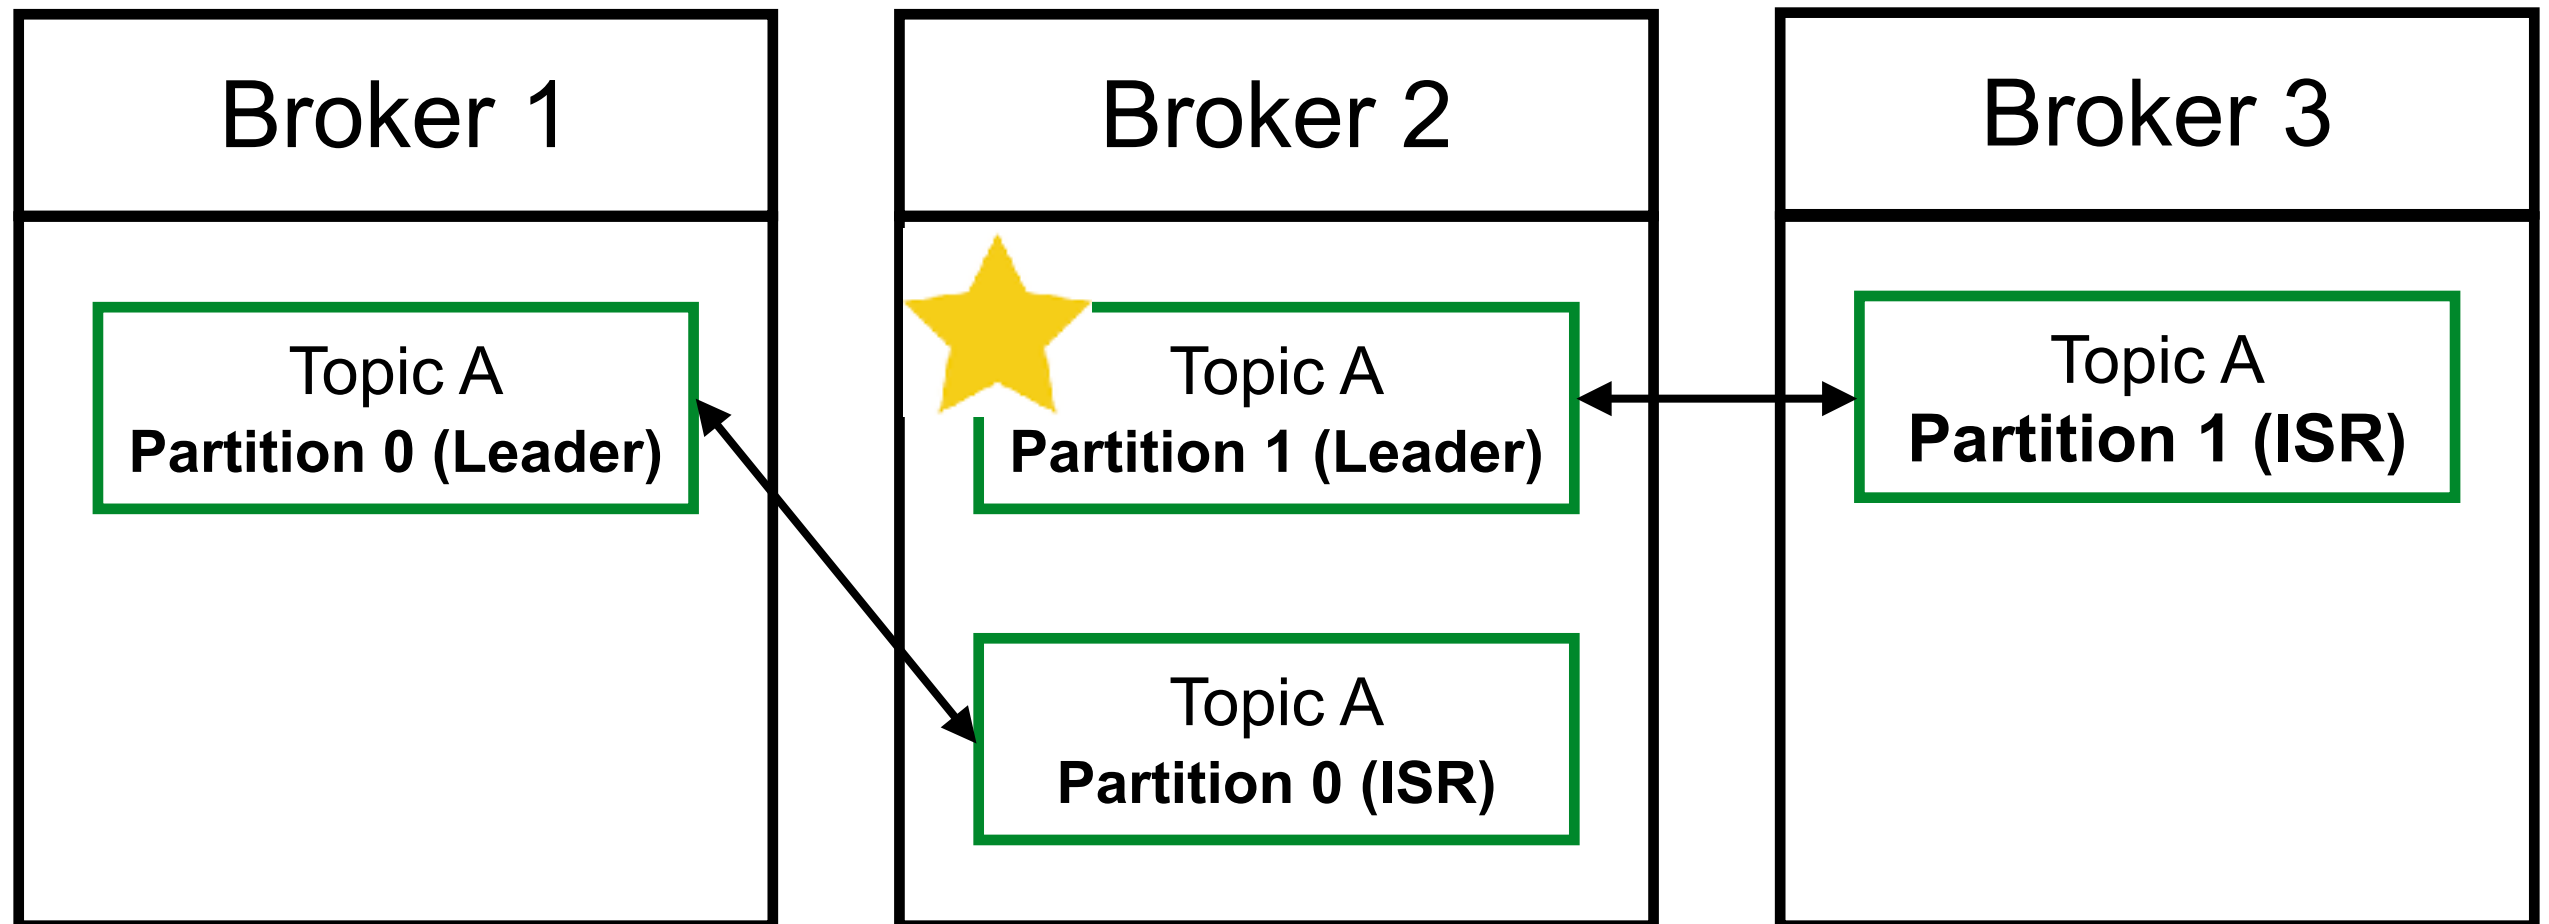
Other partition called **in-sync replica (ISR)**



# Leader for a partition

## Example

Topic A with 2 partitions and replication factor = 2



# Producers

Producers write data to topics

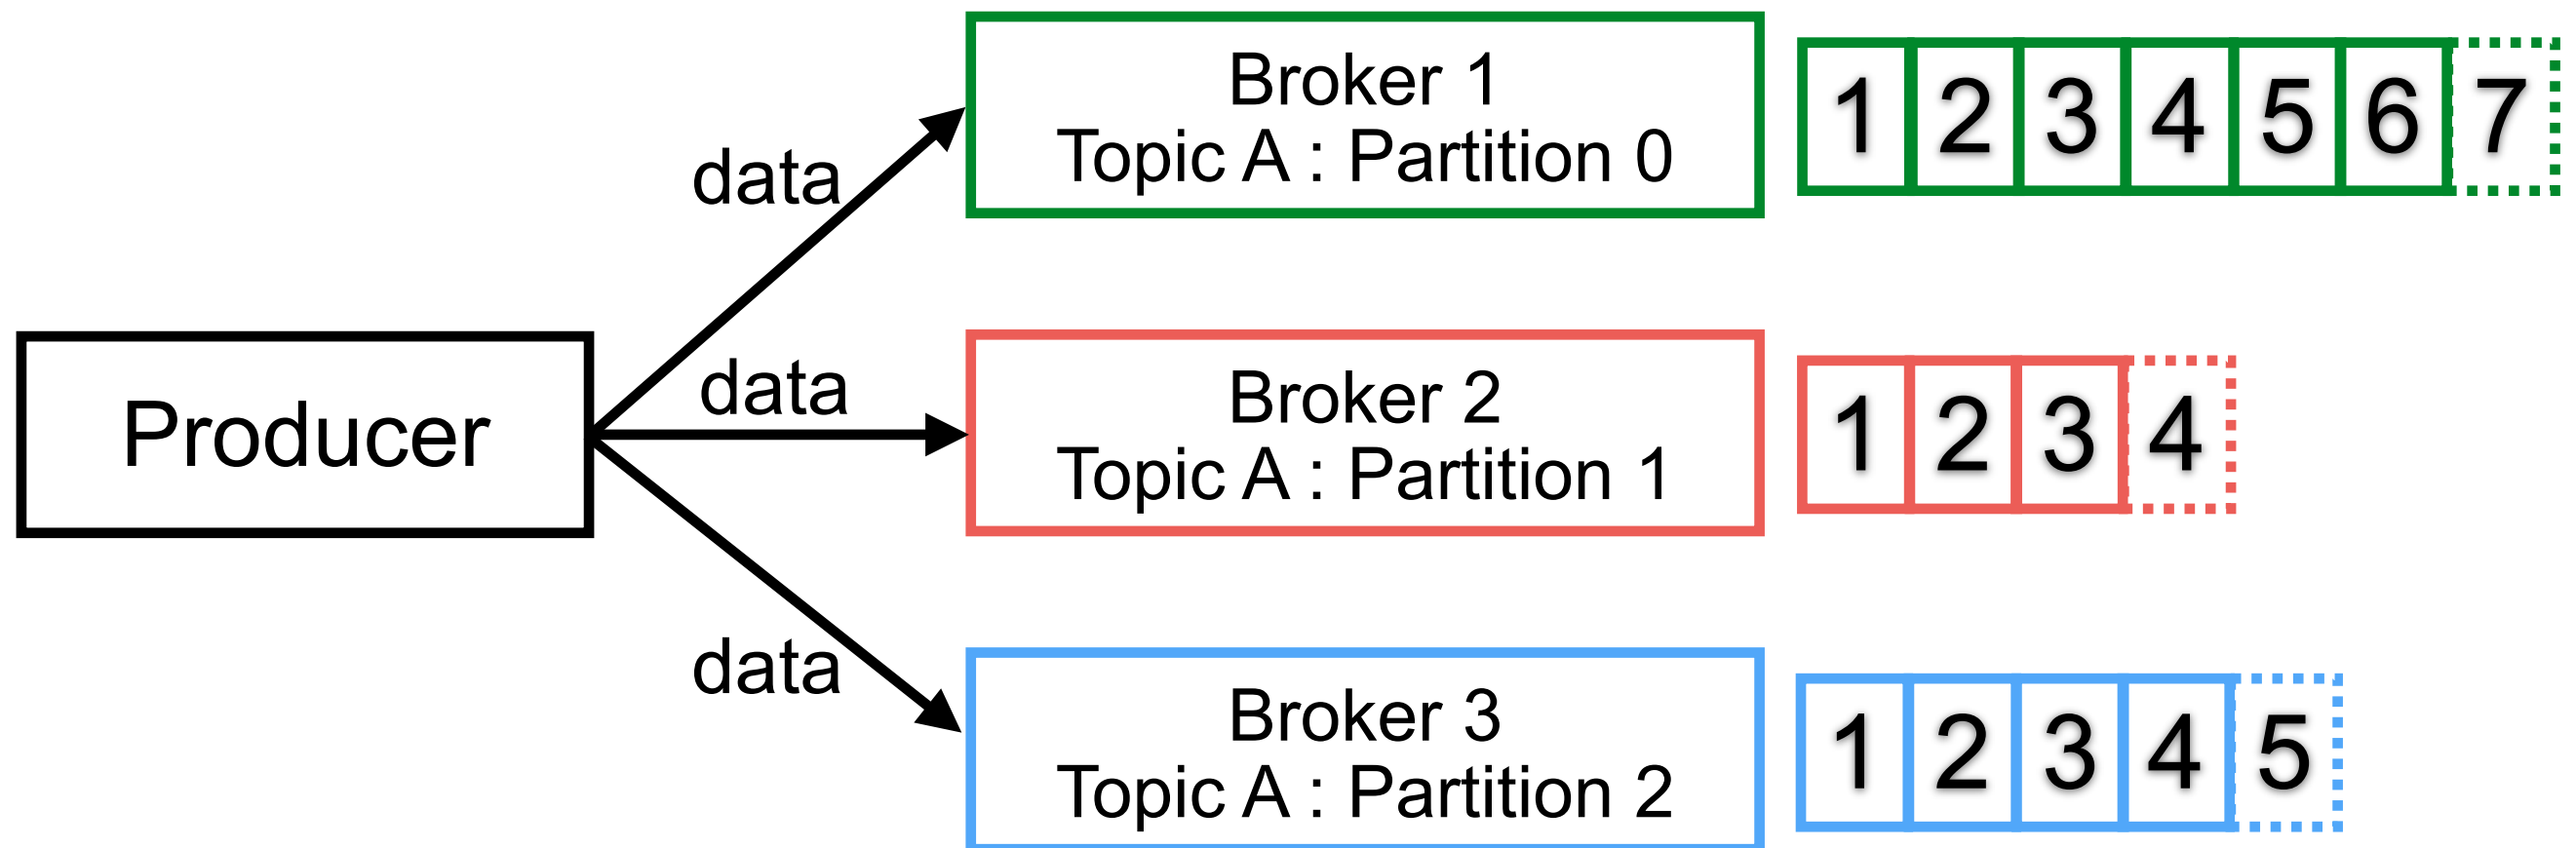
Automatically know to broker and partition to write

When broker failures, producers will automatically recover





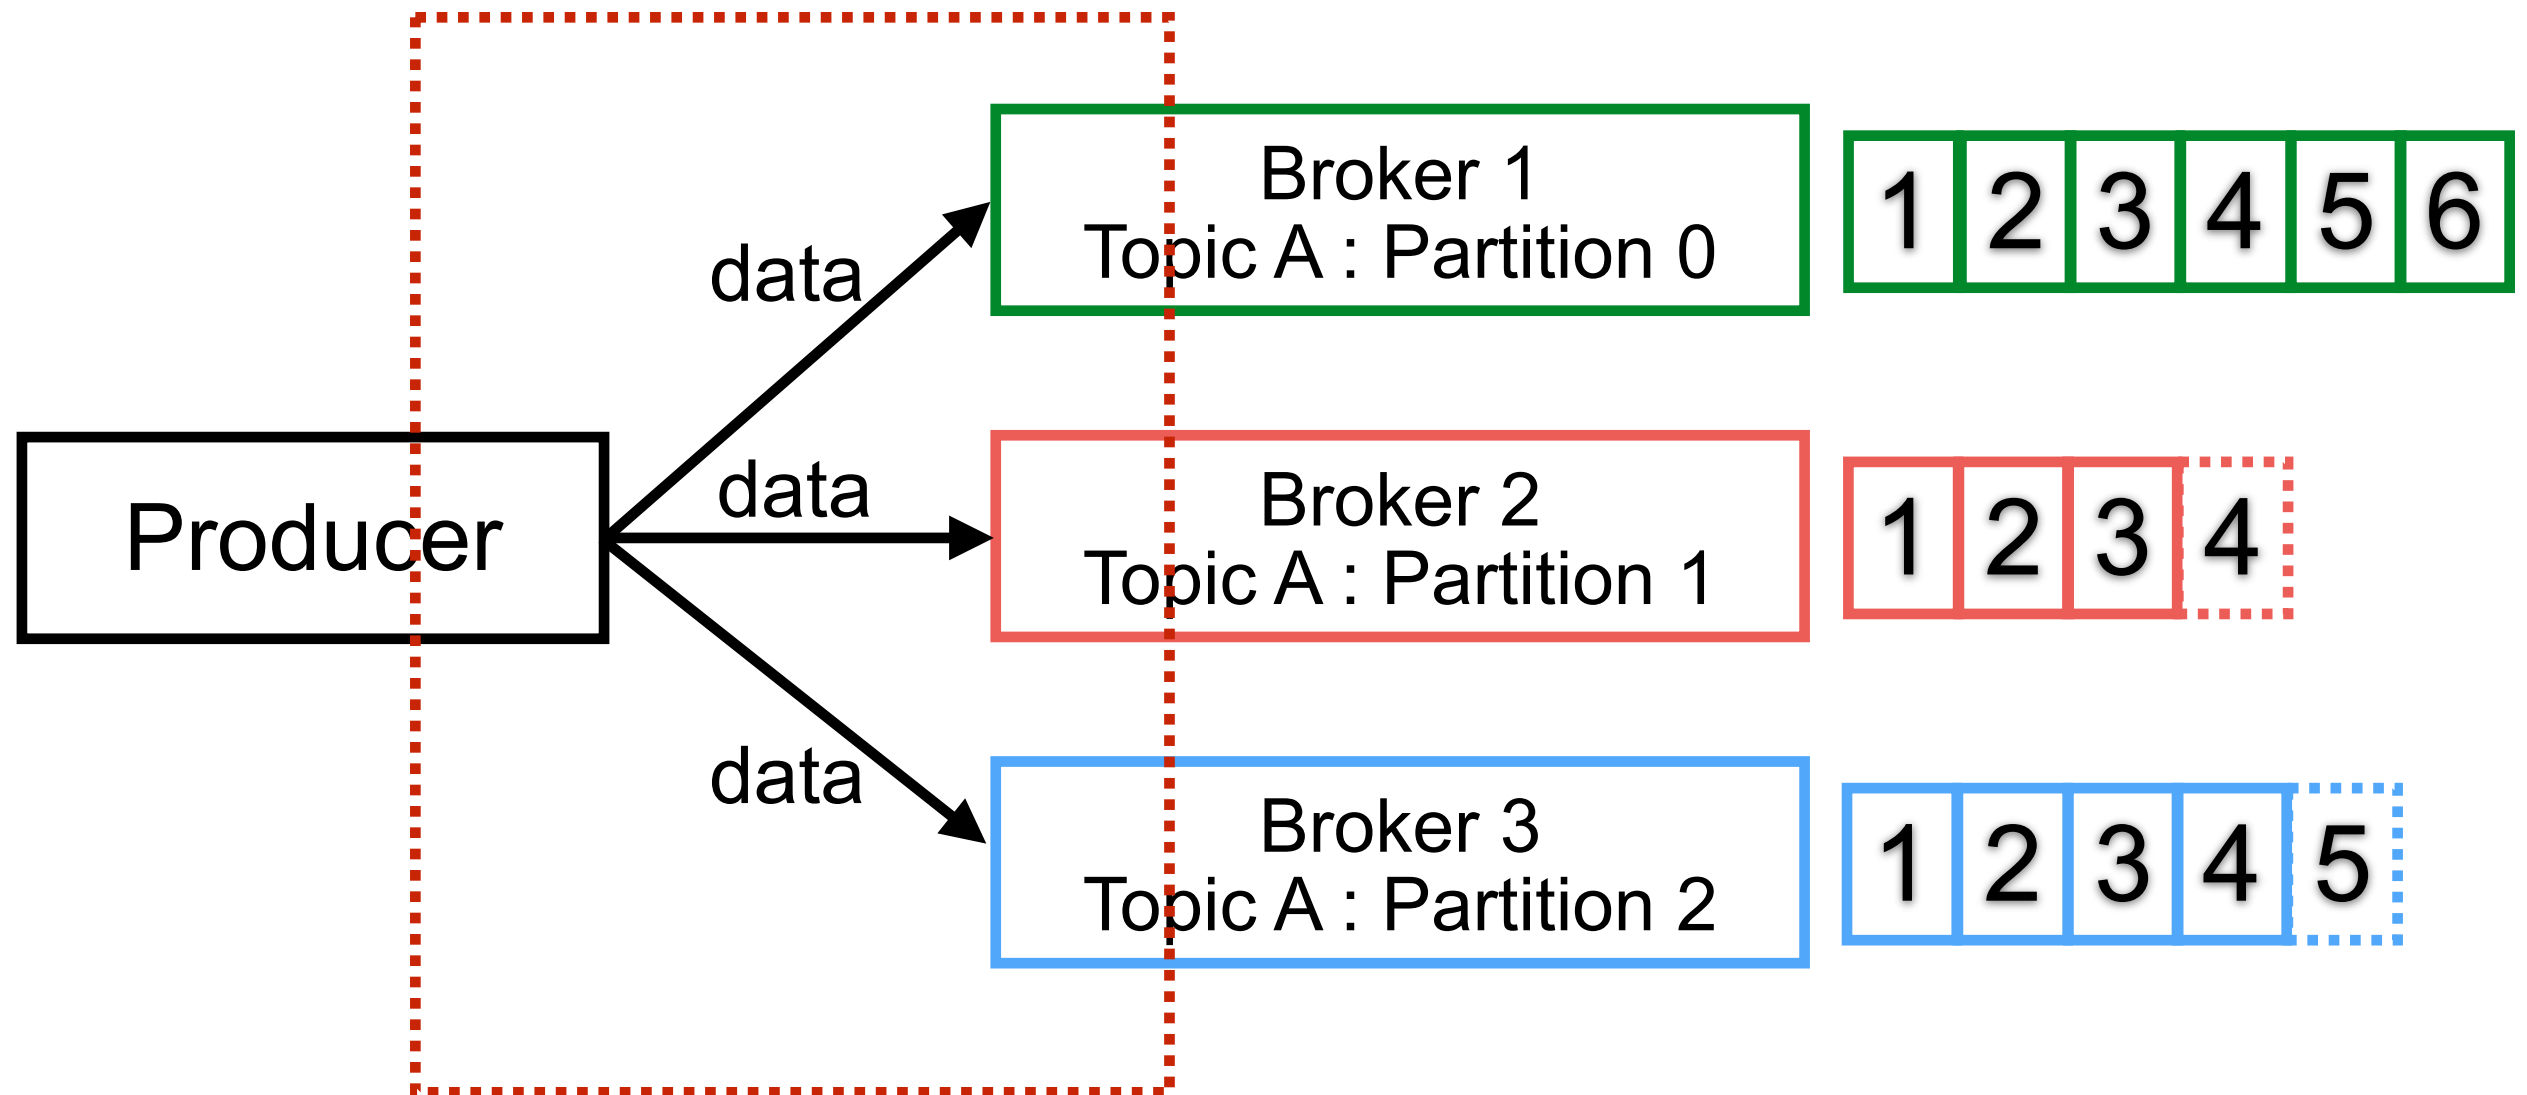
# Producers



\*\*\* *The load is balance to many brokers (no. of partitions)* \*\*\*



# Producers issue to write data !!



# Producers with acknowledgment

## **acks=0**

Producer not wait for acknowledgment

Possible data loss

## **acks=1**

Producer will wait for **leader** acknowledgment

Limited data loss

## **acks=all**

Producer will wait for **leader + ISR** acknowledgment

No data loss

<https://docs.confluent.io/platform/current/clients/producer.html>



# Partition assignment strategies

Range

Round Robin

Sticky


Customize (AbstractPartitionAssignor)

<https://docs.confluent.io/platform/current/installation/configuration/consumer-configs.html#partition-assignment-strategy>




# Producers with message keys

Producers can choose to send a **key** with data  
**Key = null**, data is sent **Sticky partitioner**



# Why use message keys ?

You need message ordering




# Consumers

Consumers read data from topic

Consumers know which broker to read from

Data will read in order **within each partition**


When broker **failures**, consumer know how to recover



# Consumer groups

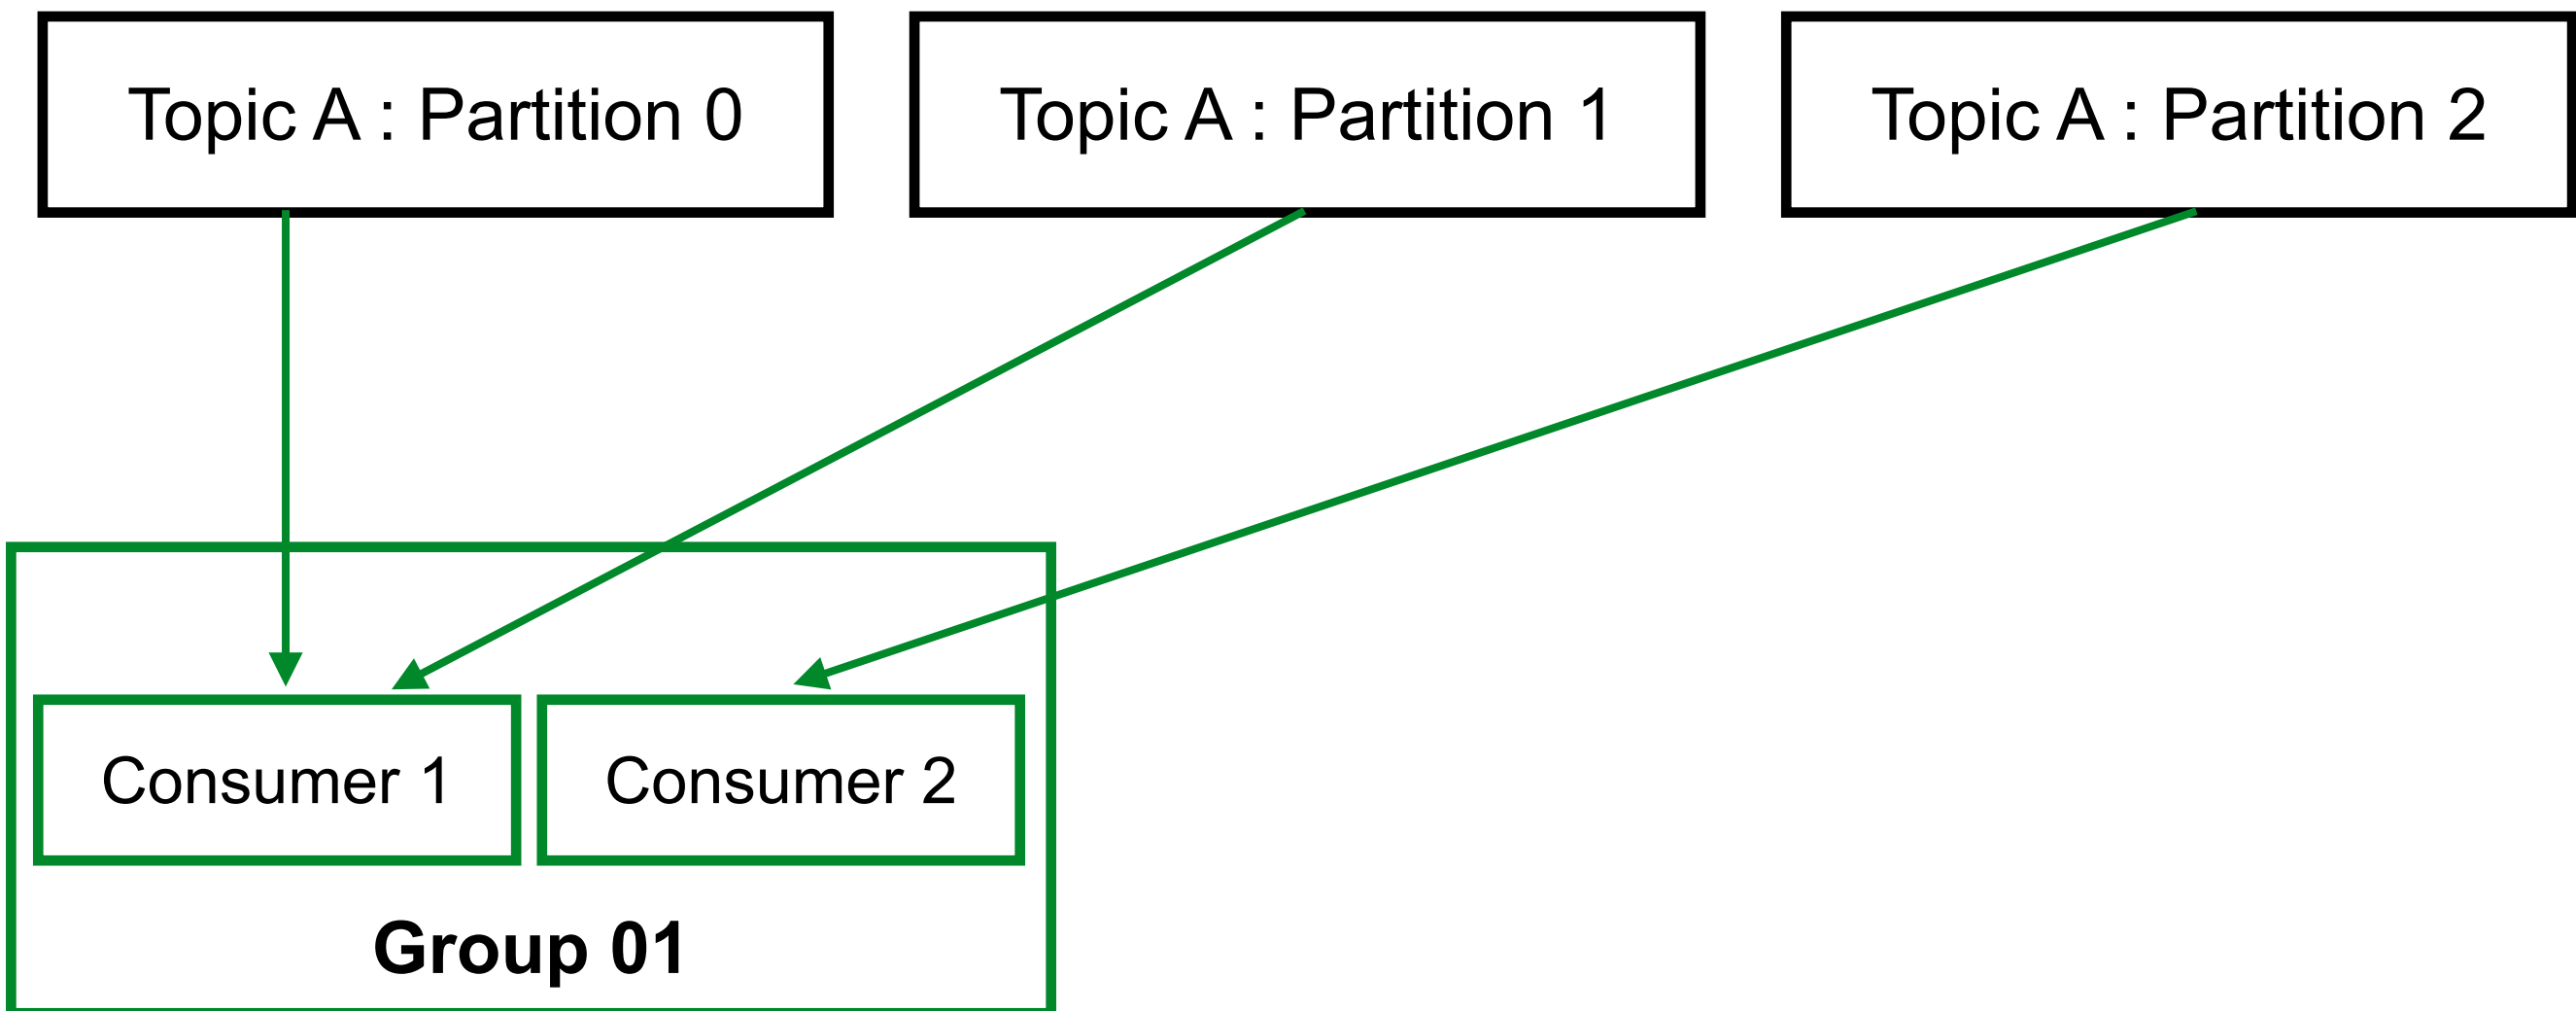
Consumers read data in consumer groups

Each consumer in a group read from exclusive partitions






# Consumer groups



Consumer will automatically use a GroupCoordinator  
ConsumerCoordinator to assign a consumer to a partition.




# Consumer groups

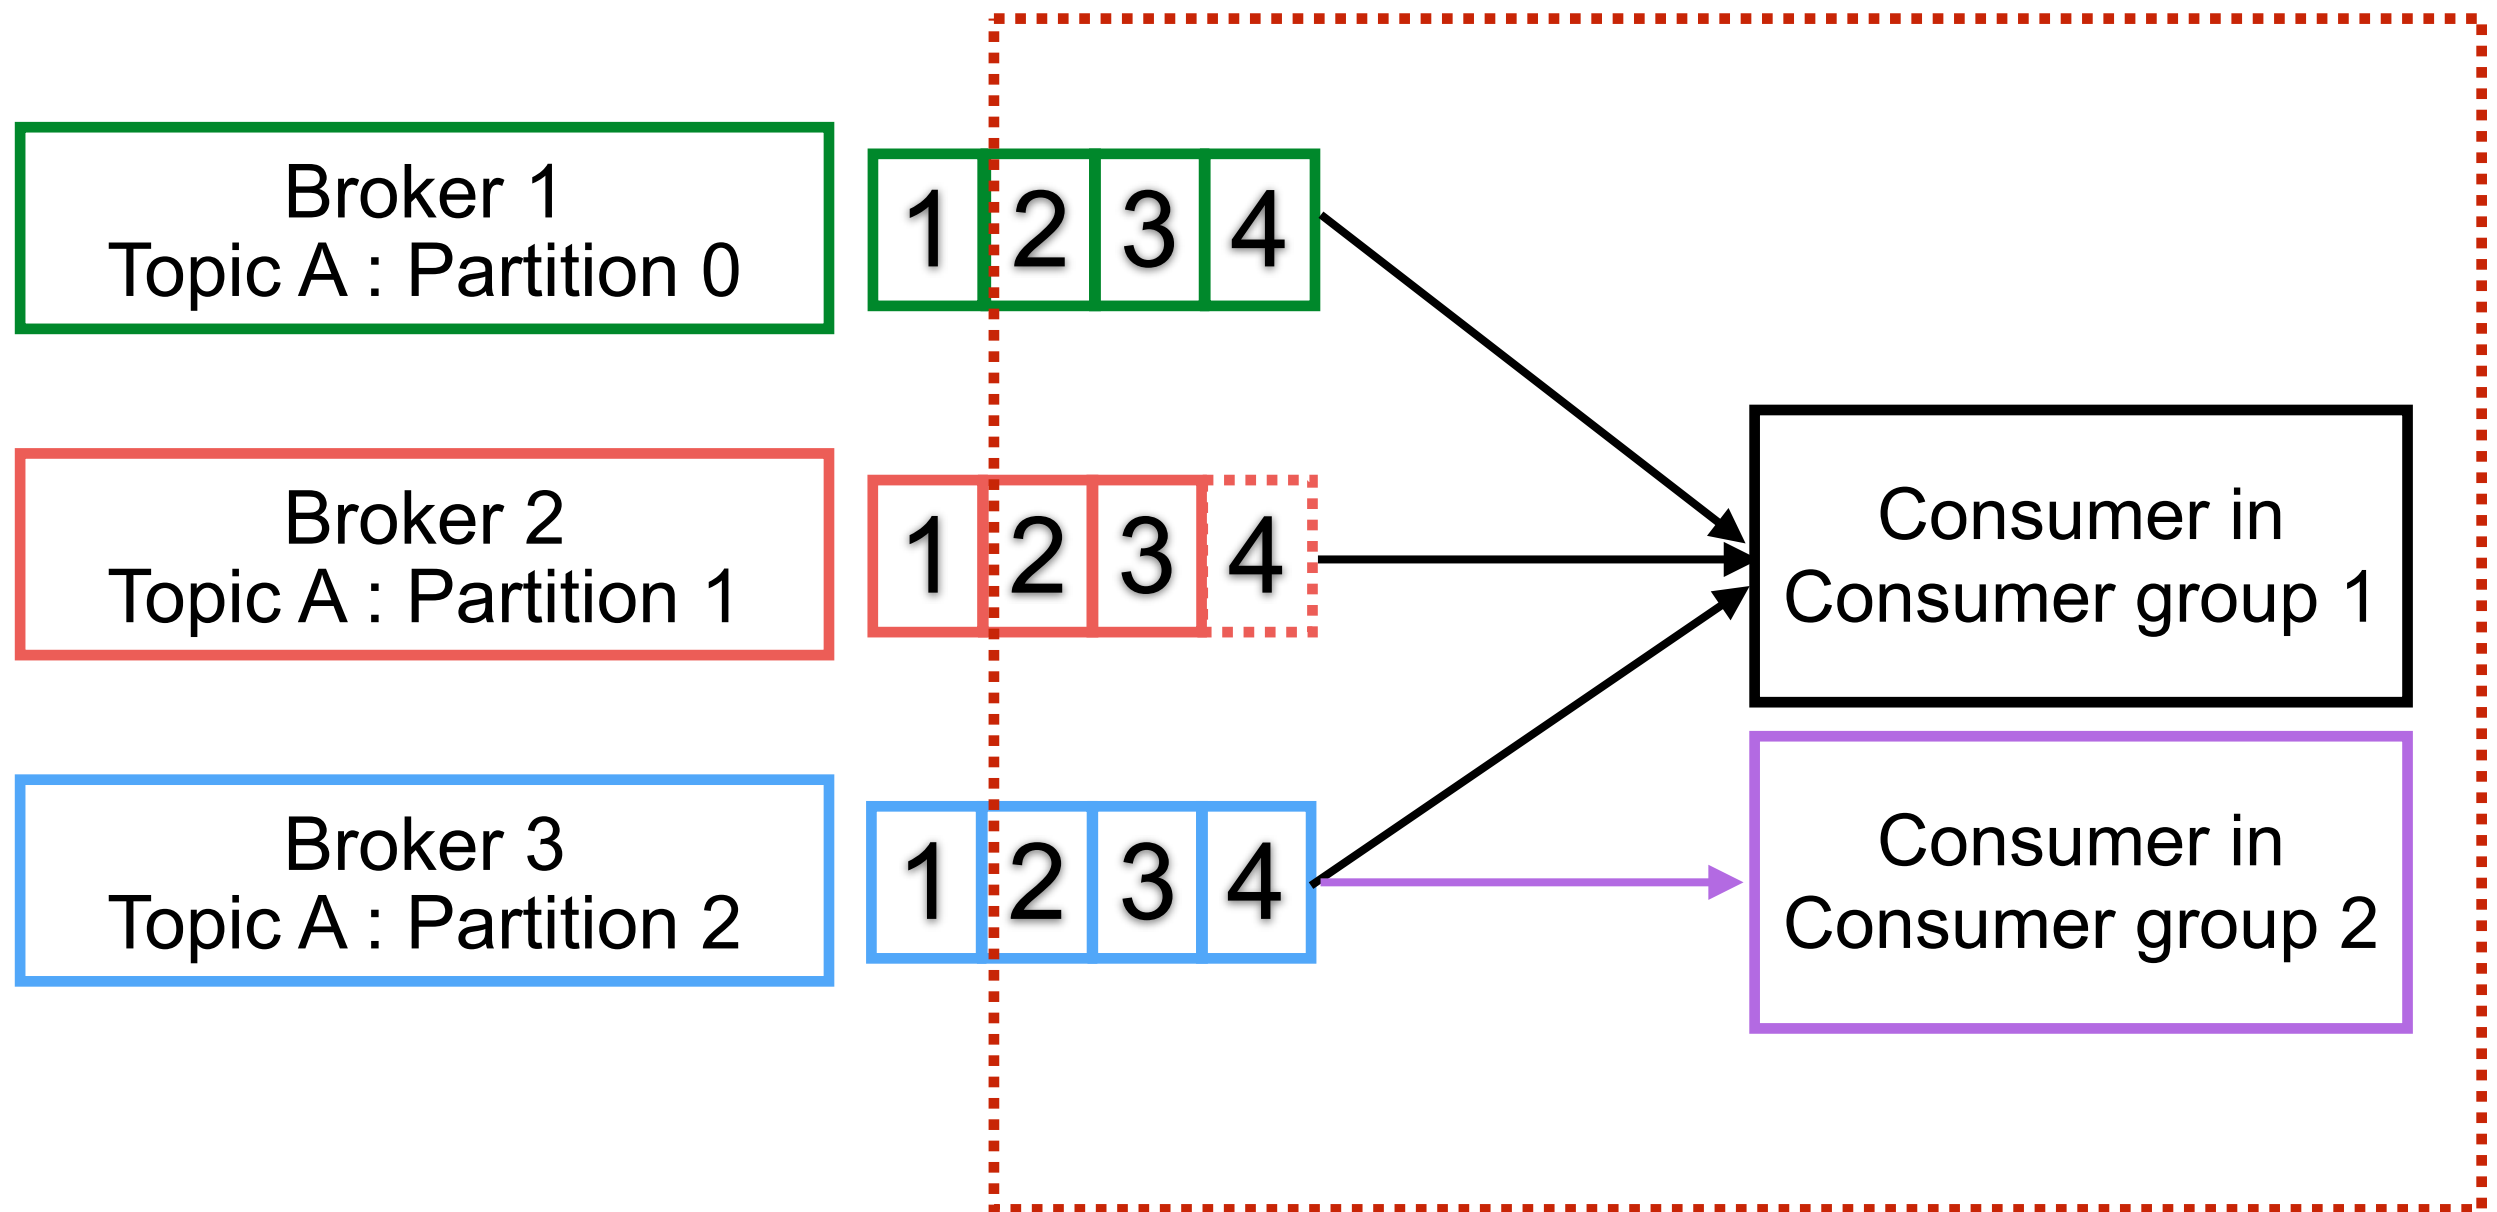


# IF Consumers > partitions ?

Some consumers will **inactive**



# Consumers offsets



# Consumers offsets

**Kafka** store the offset at which a consumer group has been reading

The offsets committed live in **topic** named  
“**\_\_consumer\_offsets**”

When consumer in a group has processed data received from Kafka,  
it should be **committing the offsets**



# When to commit the offset ?



# Delivery semantics for consumer

At most once

At least once (preferred)

Exactly once

<https://docs.confluent.io/kafka/design/delivery-semantics.html>



# 1. At most once

Offsets are committed as soon as the message is received

If processing go wrong, the message will be **loss!!**

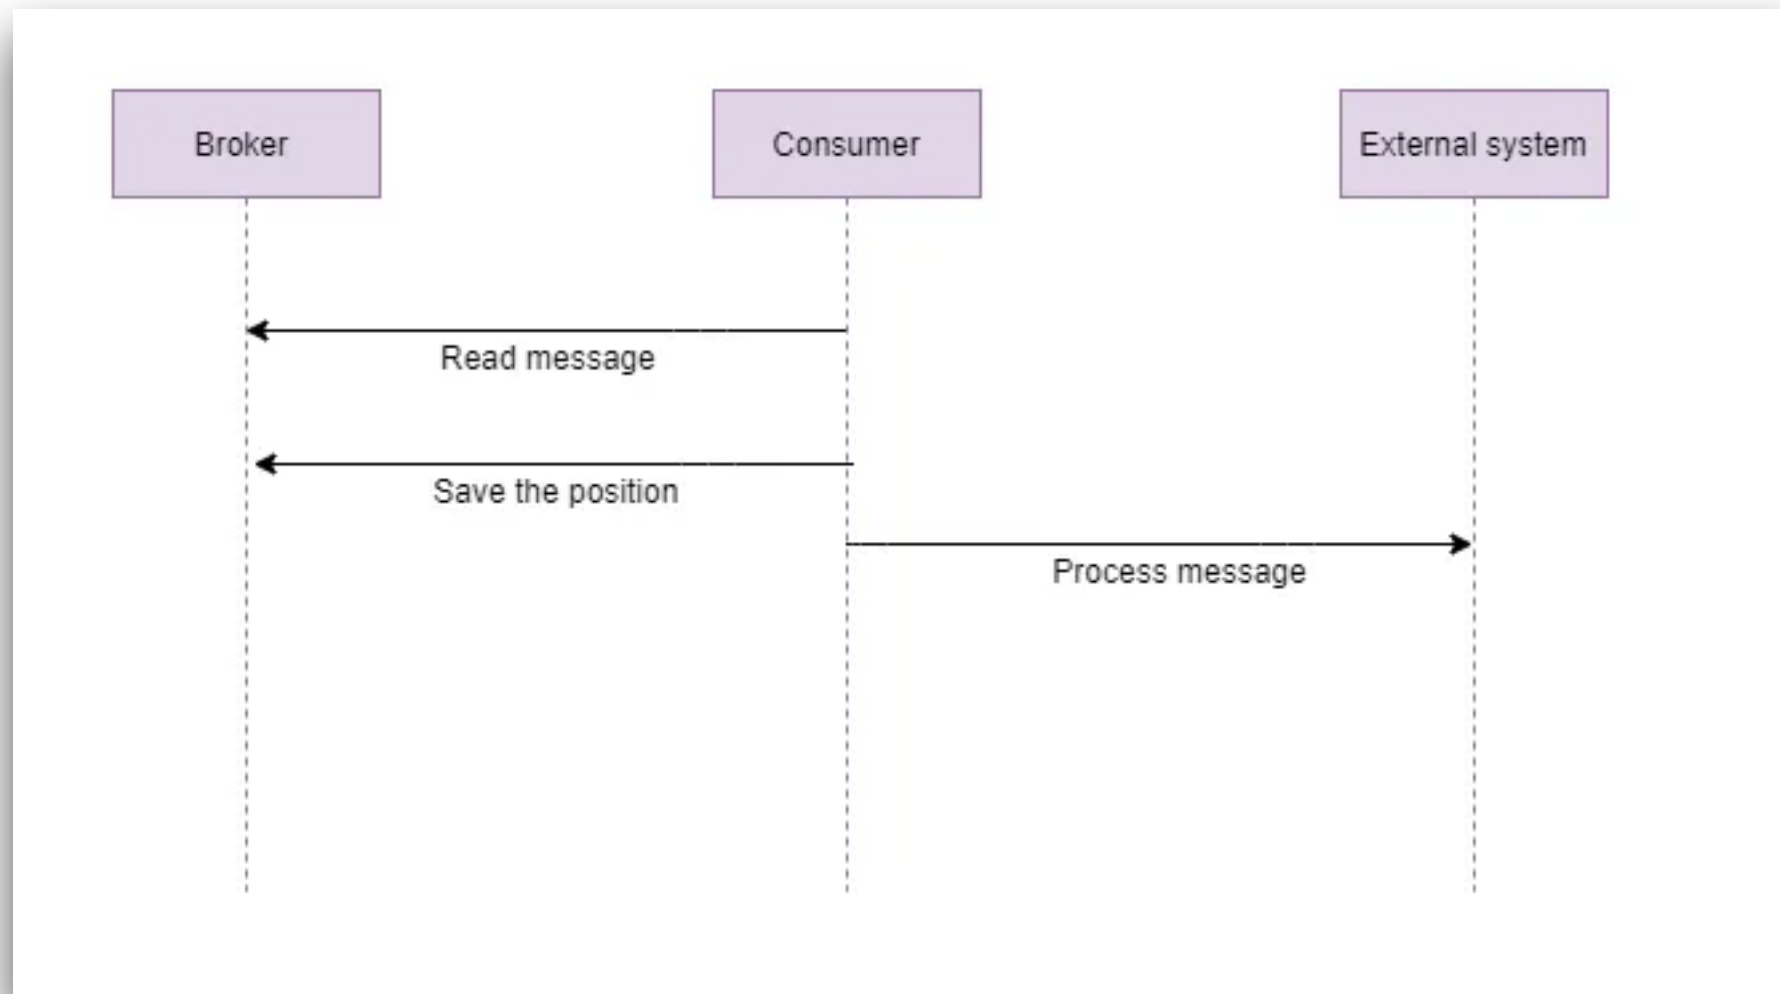
**Fire and forgot pattern**

Lowest latency





# At most once



## 2. At lease once (1)

Offsets are committed after the message is processed

Messages are never lost but may be **redelivered**

If processing go wrong, the message will be **read again**



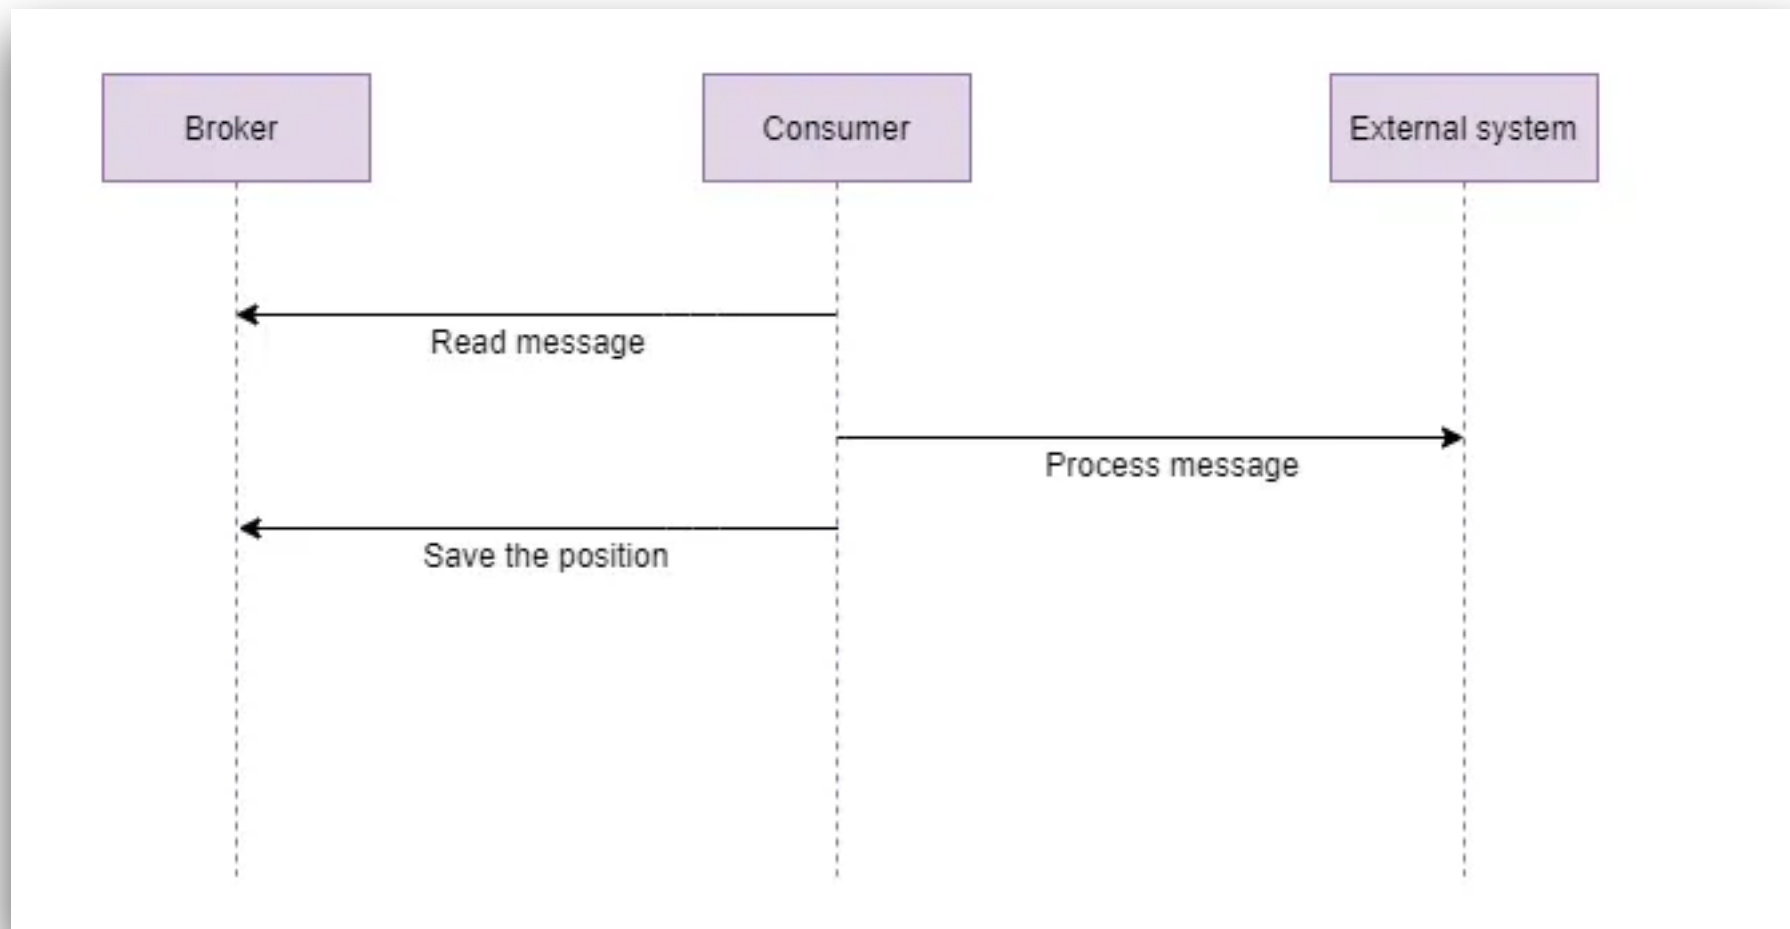
## 2. At lease once (2)

Make sure your processing is **idempotent**

*Processing again the message  
not impact to your system !!*



# At lease once



# 3. Exactly once

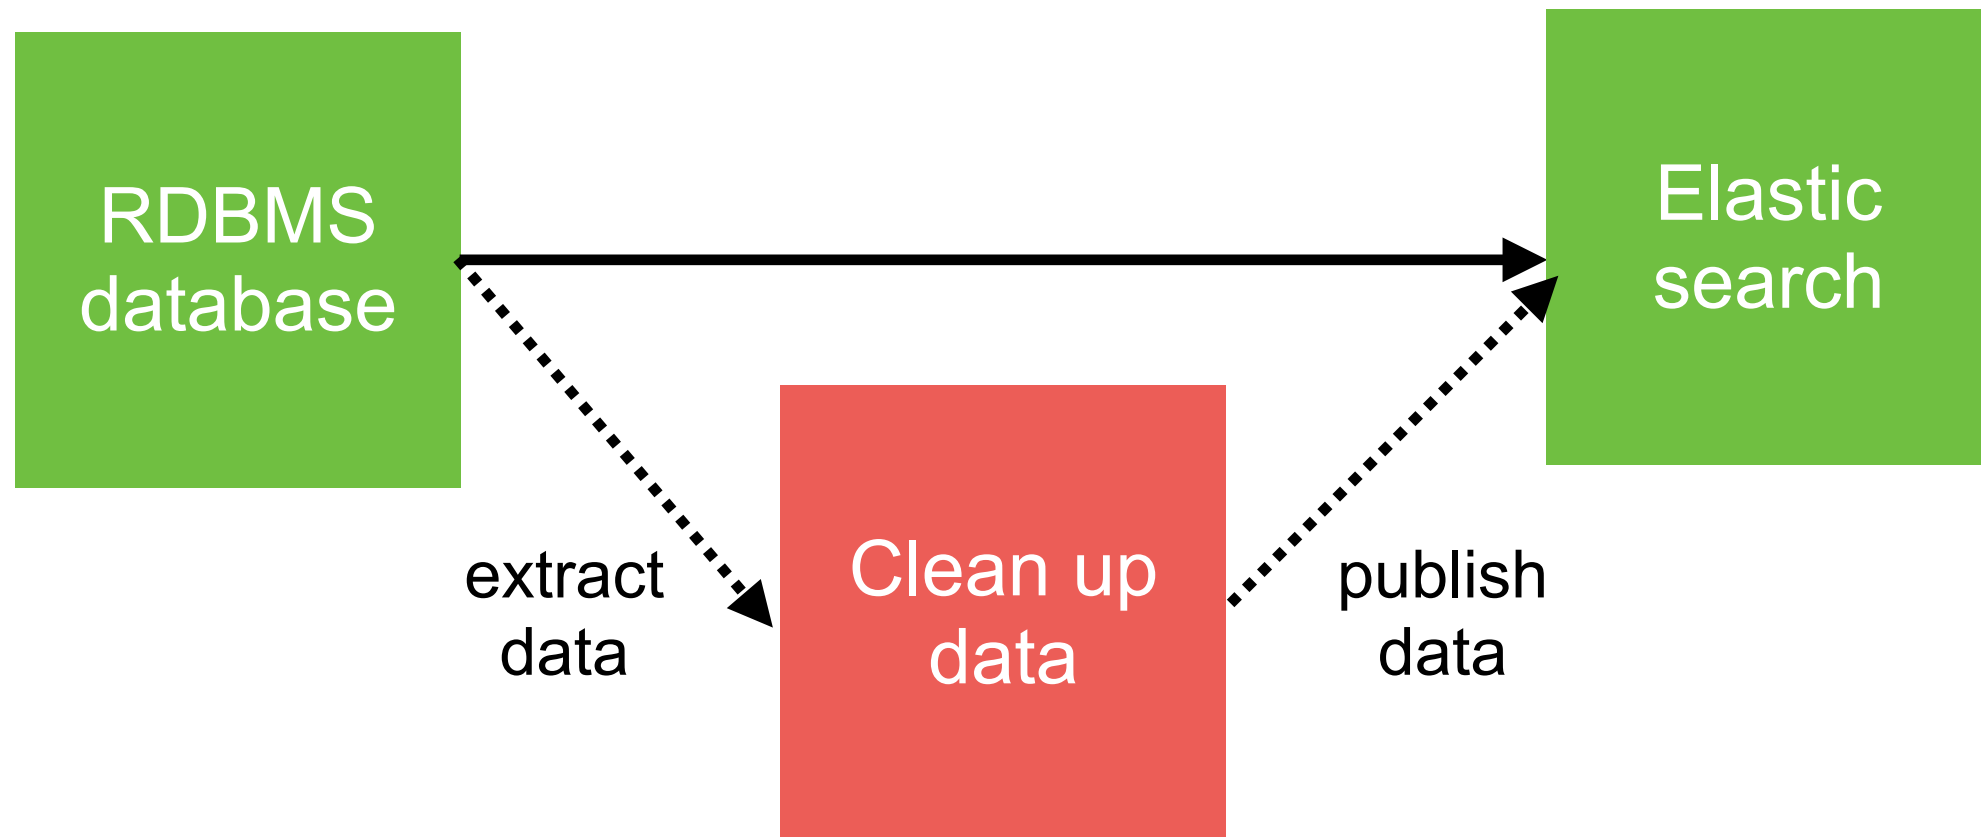
Each message is delivered once and only once  
Can be achieved for Kafka (Workflow, Stream API)  
Transactional delivery  
Resent with **idempotent**



# Data Pipeline




# Batching process



Schedule task run every 2 am



# Streaming data pipeline




Real-time(NRT) process when data changed






# Process !!



<https://dataengineeringcentral.substack.com/p/batch-vs-near-realtime-vs-streaming>



# Example



<https://github.com/up1/workshop-kafka-2025/tree/main/kafka-cdc>



# Q/A

